

# Nearly Unfathomable Morass of Arcana (NUMA)


Fedor G Pikus  
Technical Fellow

# Non-Uniform Memory Architecture (NUMA)


Fedor G Pikus  
Technical Fellow

# Short version – slide 1: This is NUMA


Computers went from this ...



... to this



# Short version – slide 2: This happens if you ignore NUMA



# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

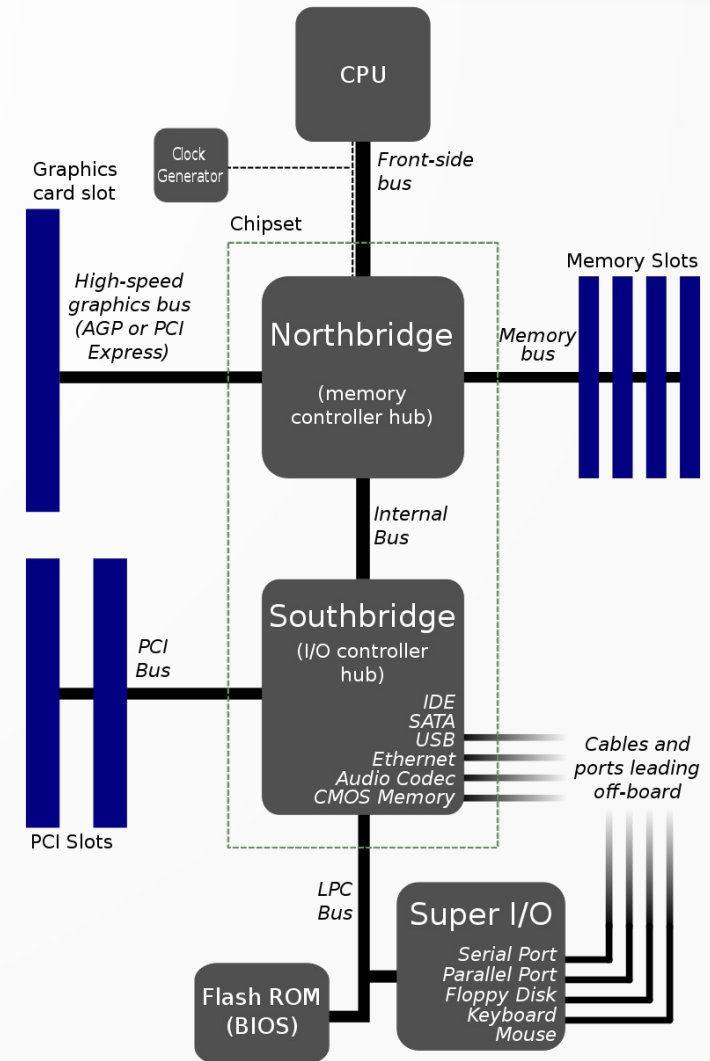



# What is NUMA

- Non-Uniform Memory Architecture
- What is Uniform Memory Architecture?

# Uniform Memory Architecture

- All processors have the same access to all memory
  - X86 Front-Side Bus (FSB)



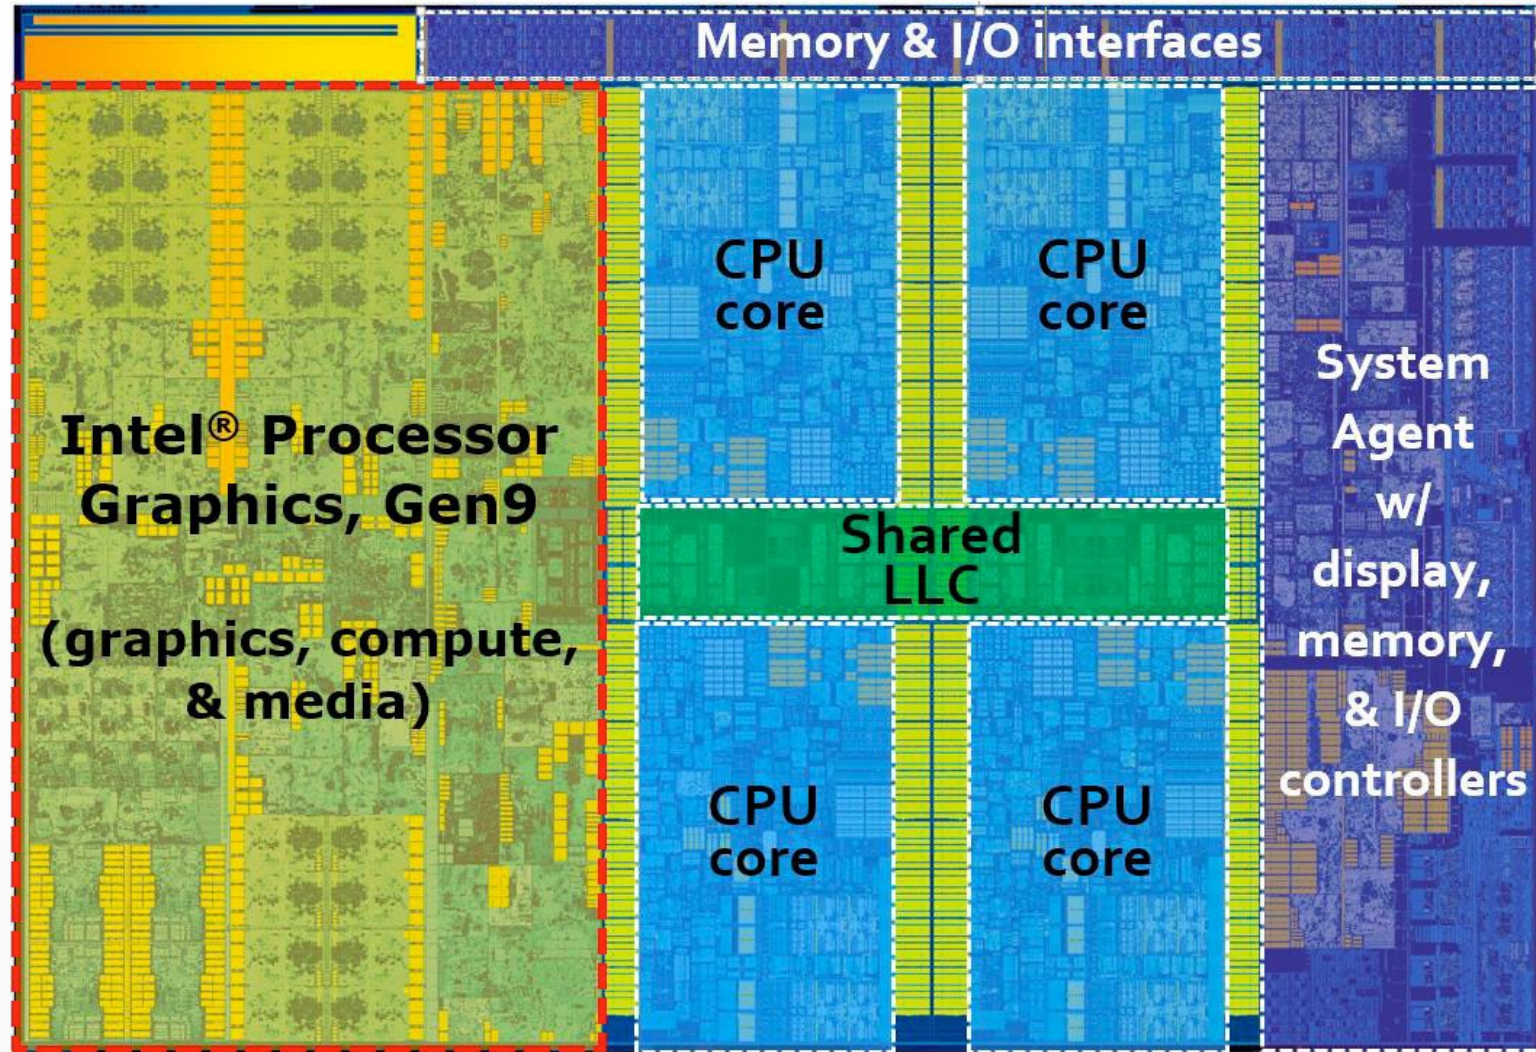


# Uniform Memory Architecture

- UMA is not just for ancient hardware

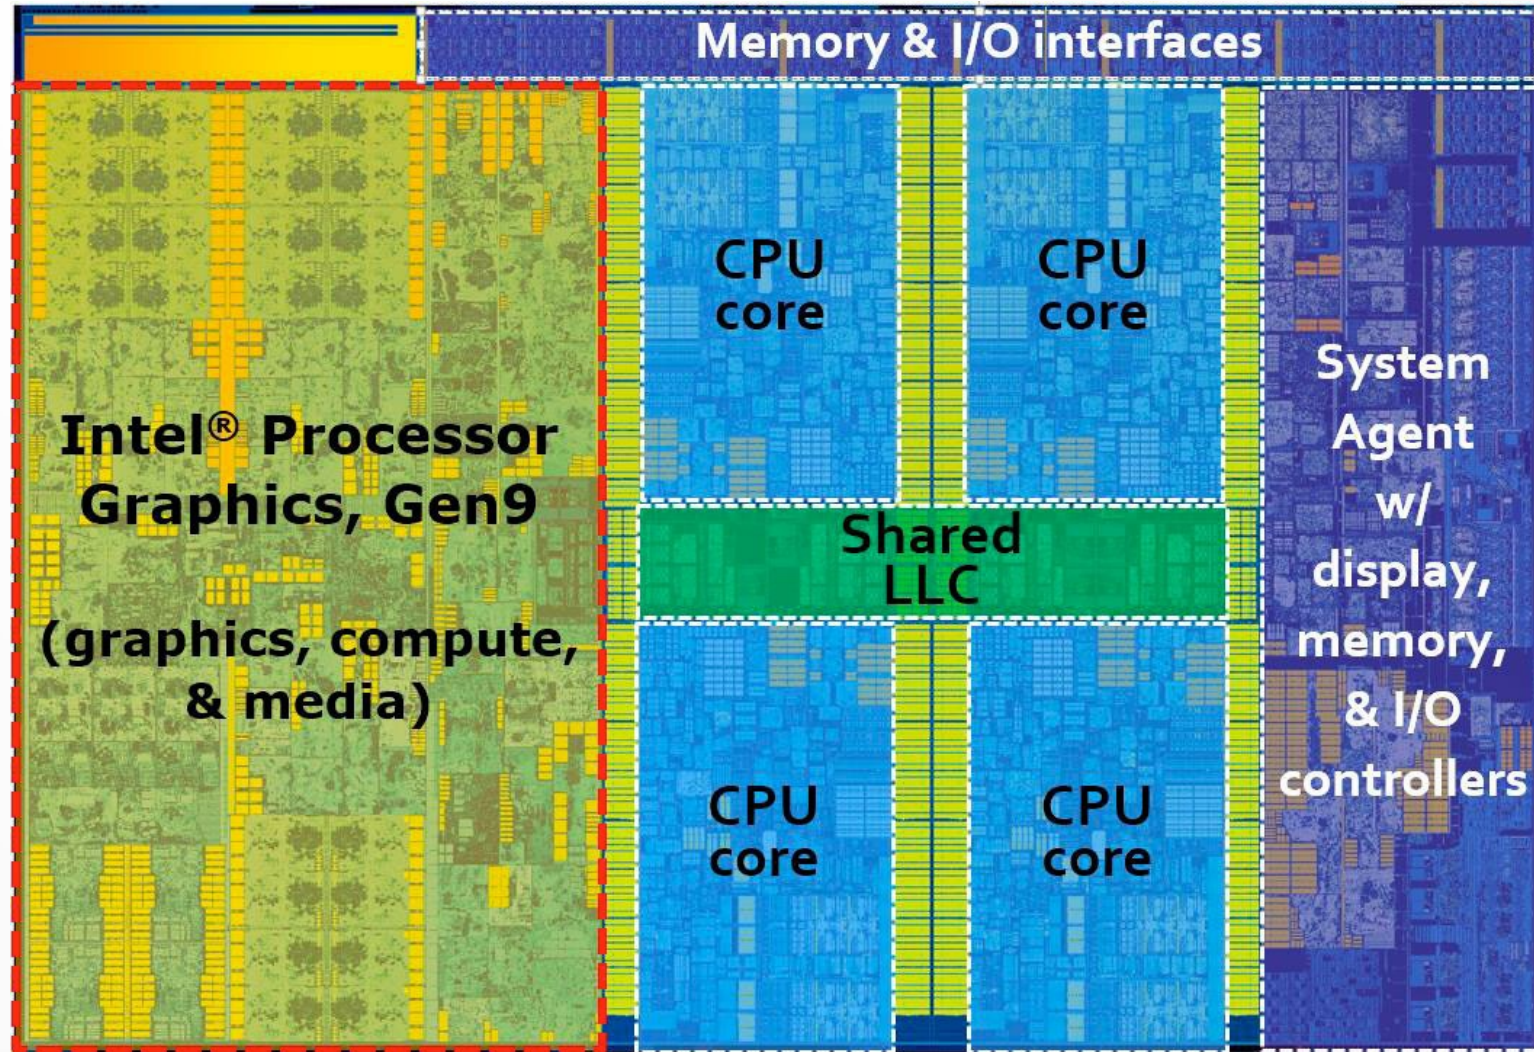

# Uniform Memory Architecture

- UMA is not just for ancient hardware
- Intel SkyLake



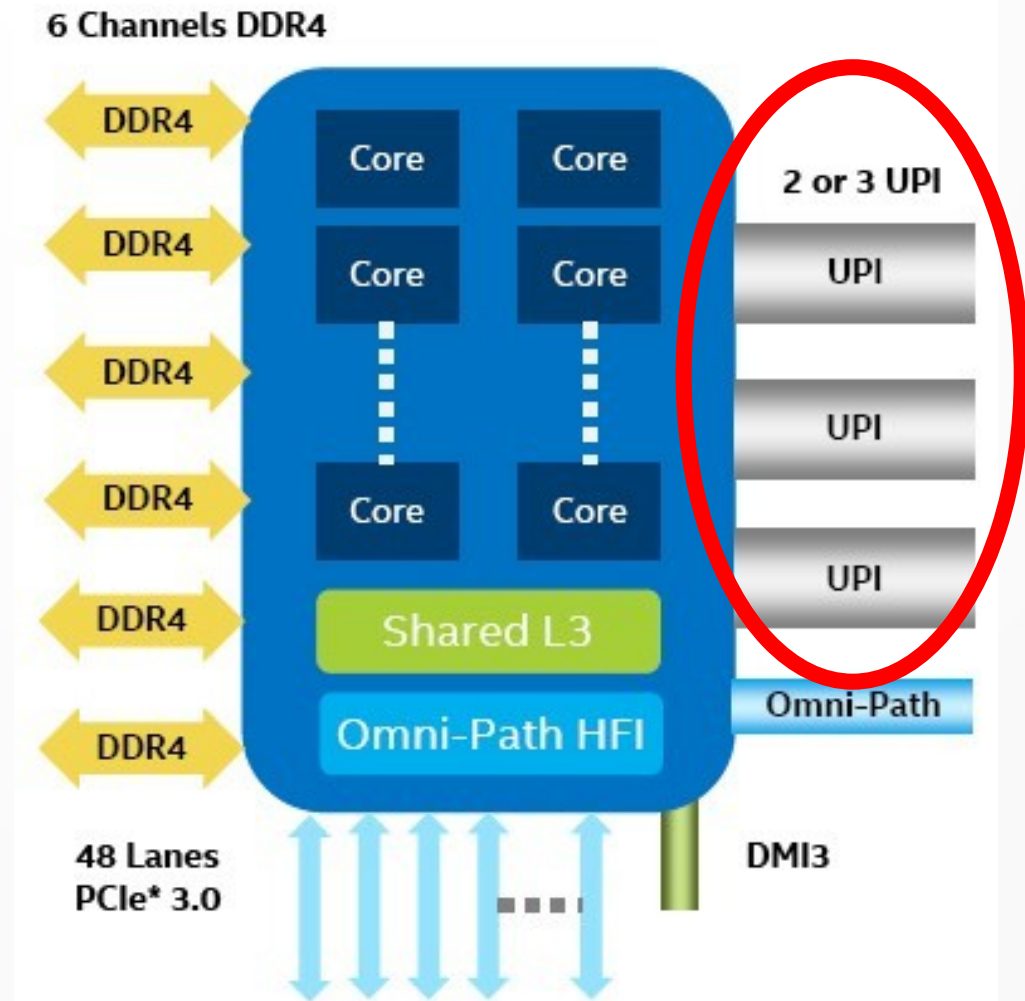


# Uniform Memory Architecture



# Non-Uniform Memory Architecture

- Aren't all systems with modern processors NUMA?
  - They can be...
- Also Intel SkyLake →

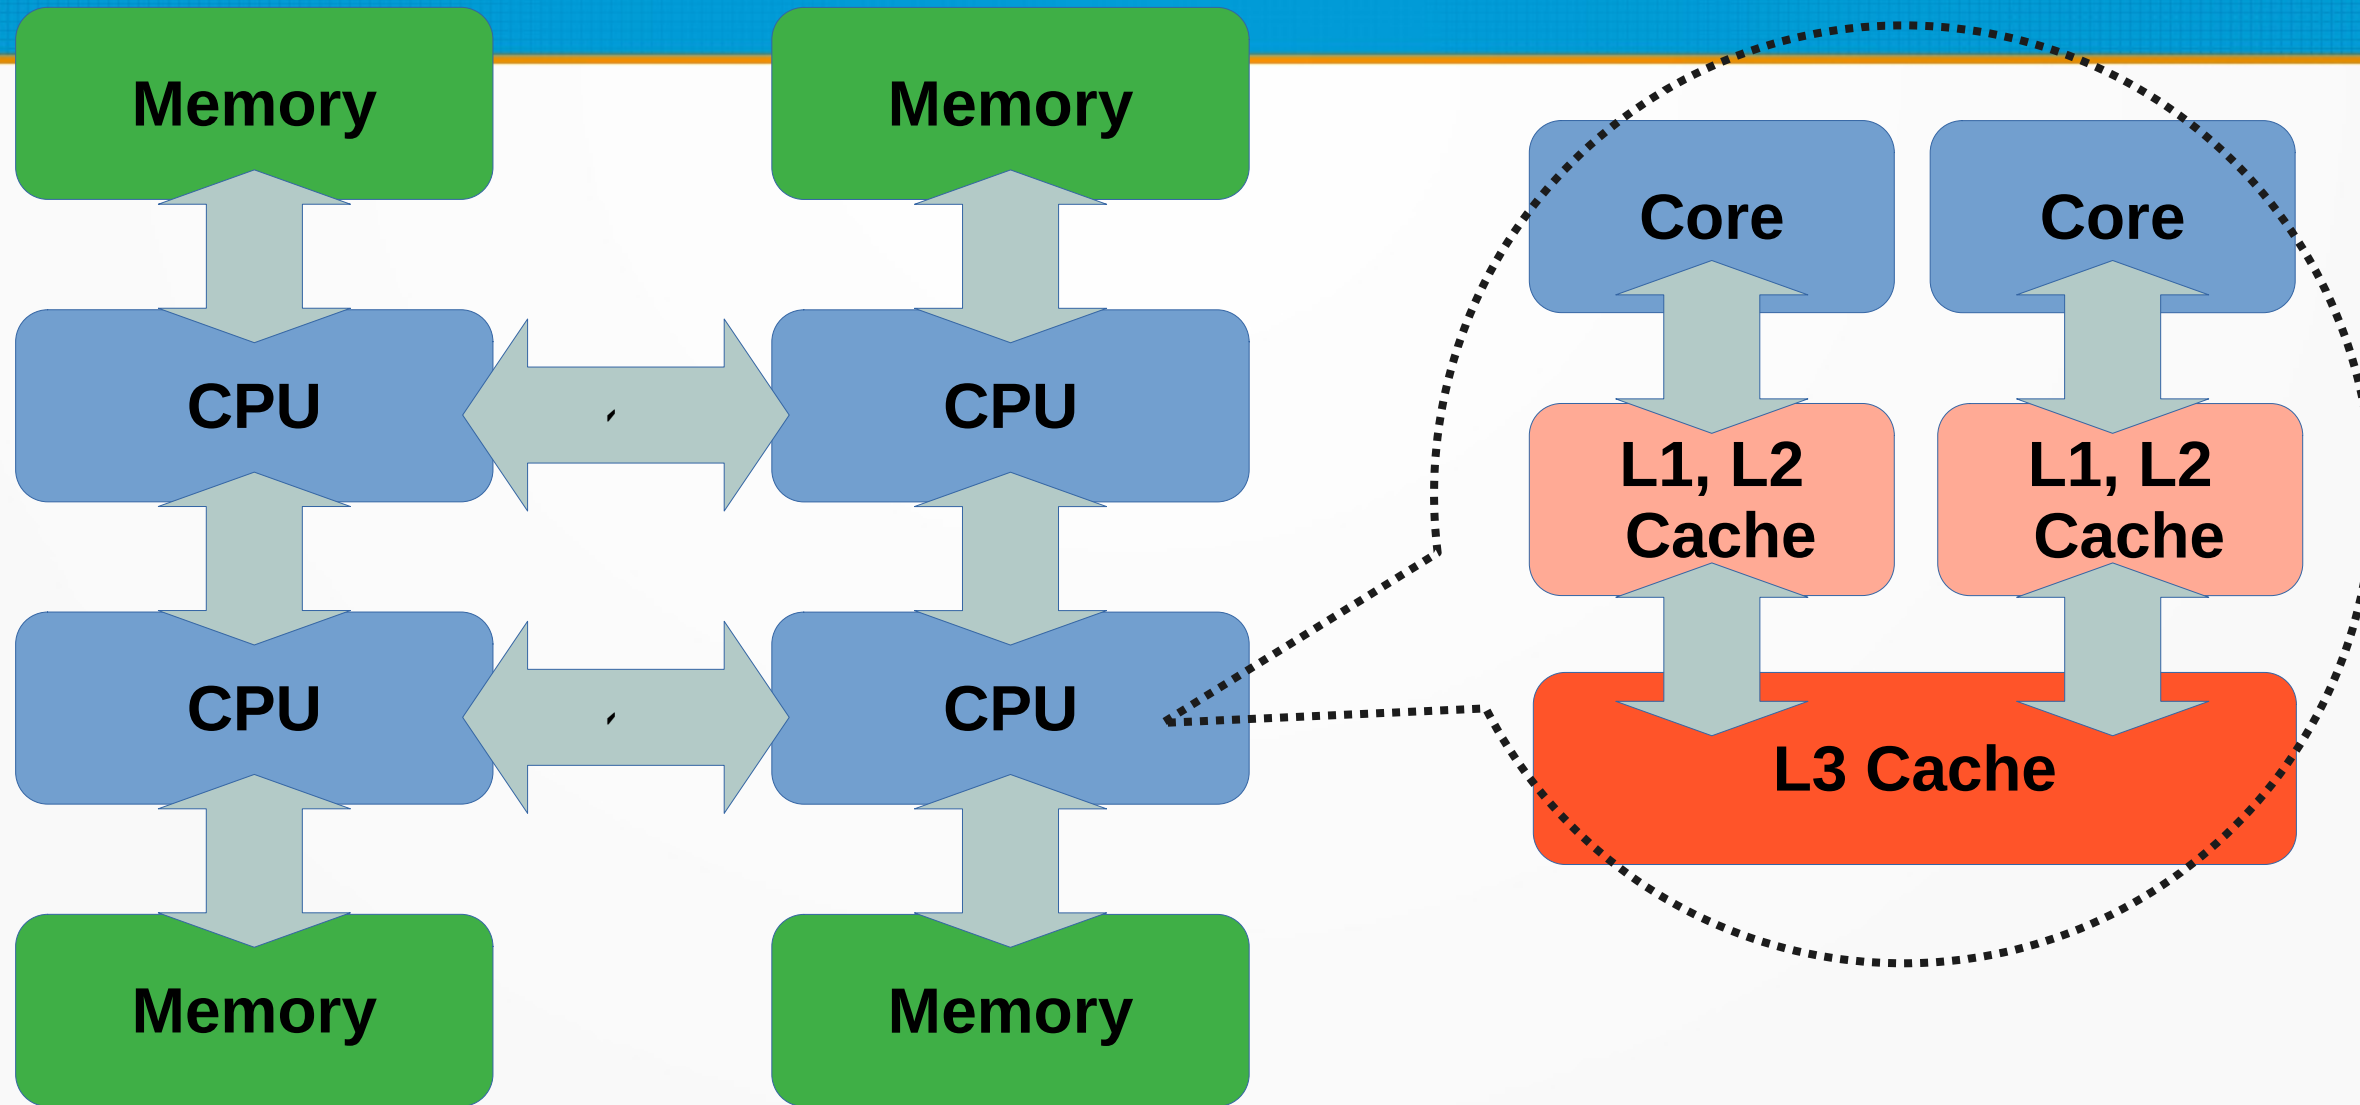




# What is NUMA

- Non-Uniform Memory Architecture
- What is Uniform Memory Architecture?
- What is Non-Uniform Memory Architecture?


# Non-Uniform Memory Architecture





# Brief history of NUMA


- AMD Opteron (2003) – Hyperconnect
- Intel Nehalem (2007) – Quick Path Interconnect (QPI)
  - Intel SkyLake (2017) – UltraPath Interconnect (UPI)
- X86 systems with integrated memory controllers
  - Replaces chipset controllers (North and South Bridge)





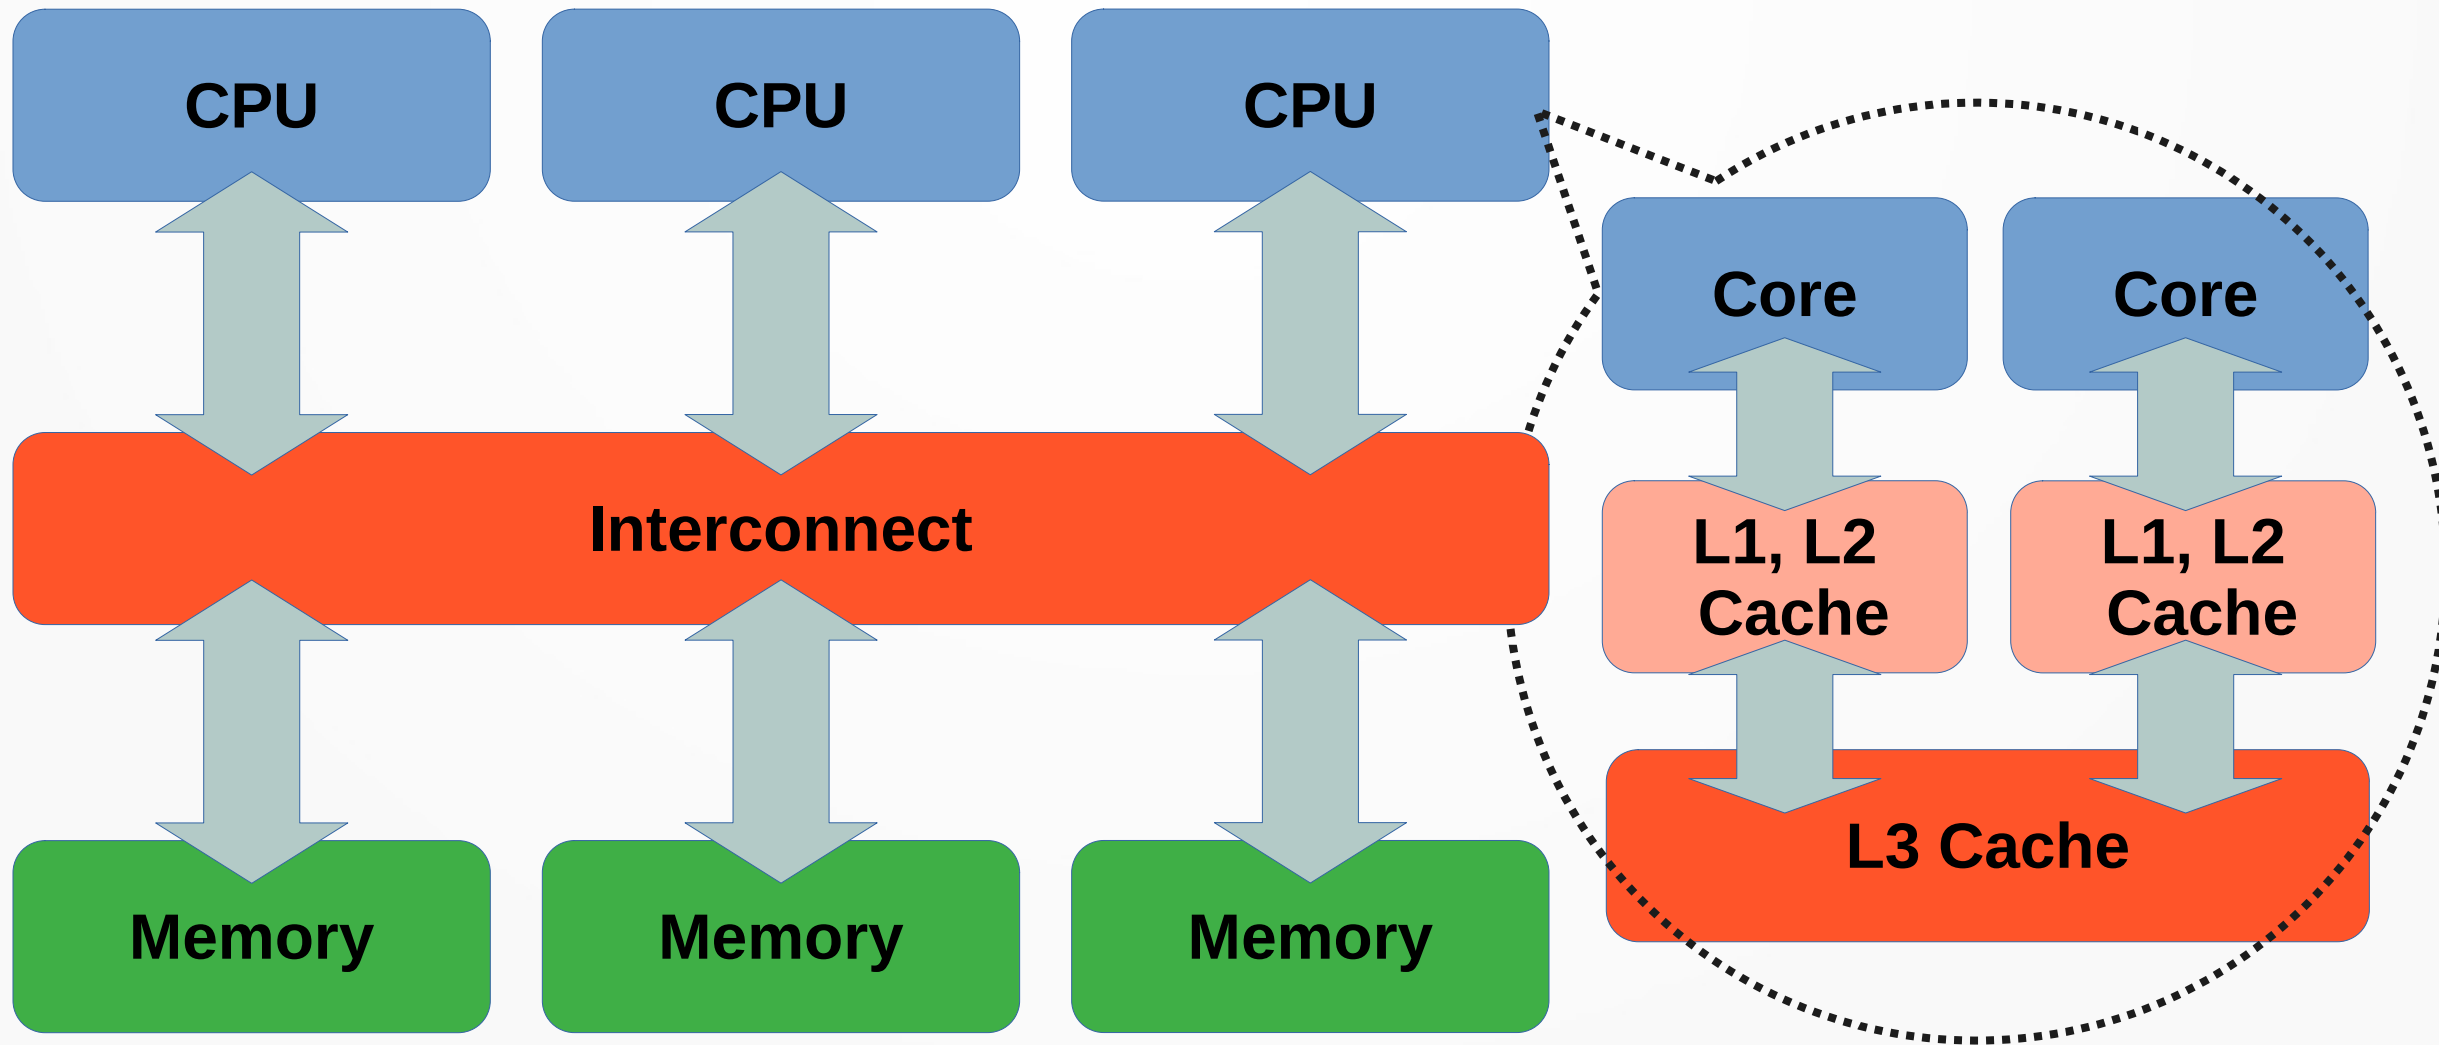
# Brief history of NUMA

- 199? - Honeywell, HP, Cray/SGI, etc
- AMD Opteron (2003) – Hyperconnect
- Intel Nehalem (2007) – QPI





# Not all multi-socket systems are NUMA




# Why NUMA? Why now?

- Systems with multiple CPUs became common
  - Not just for users with unlimited \$\$\$
- High-speed symmetric interconnect is expensive
  - Not just in \$\$\$\$s but in power, heat, and complexity
- On-chip memory controllers favor NUMA
  - Each CPU accesses its own memory
- Many large systems avoided the issue by using explicit message passing
  - No global address space – no NUMA
- Not just exotic systems anymore
  - Intel E5-2640 (Haswell) 16 cores \* 2 sockets – 2014, under \$1000 per CPU

# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# Non-Uniform Memory Architecture




- Process or thread runs on a specific CPU at any time
  - In time, processes can move between CPUs and nodes
- Process or thread accesses memory that resides on a particular node
  - Could be the same node or different




# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is different from accessing non-local memory
- To access memory from another node, that node's CPU must be involved


# Memory Bandwidth in NUMA systems




# Memory Bandwidth in NUMA systems



# Memory Bandwidth in NUMA systems



# Memory Bandwidth in (newer) NUMA systems



# We Will Learn:

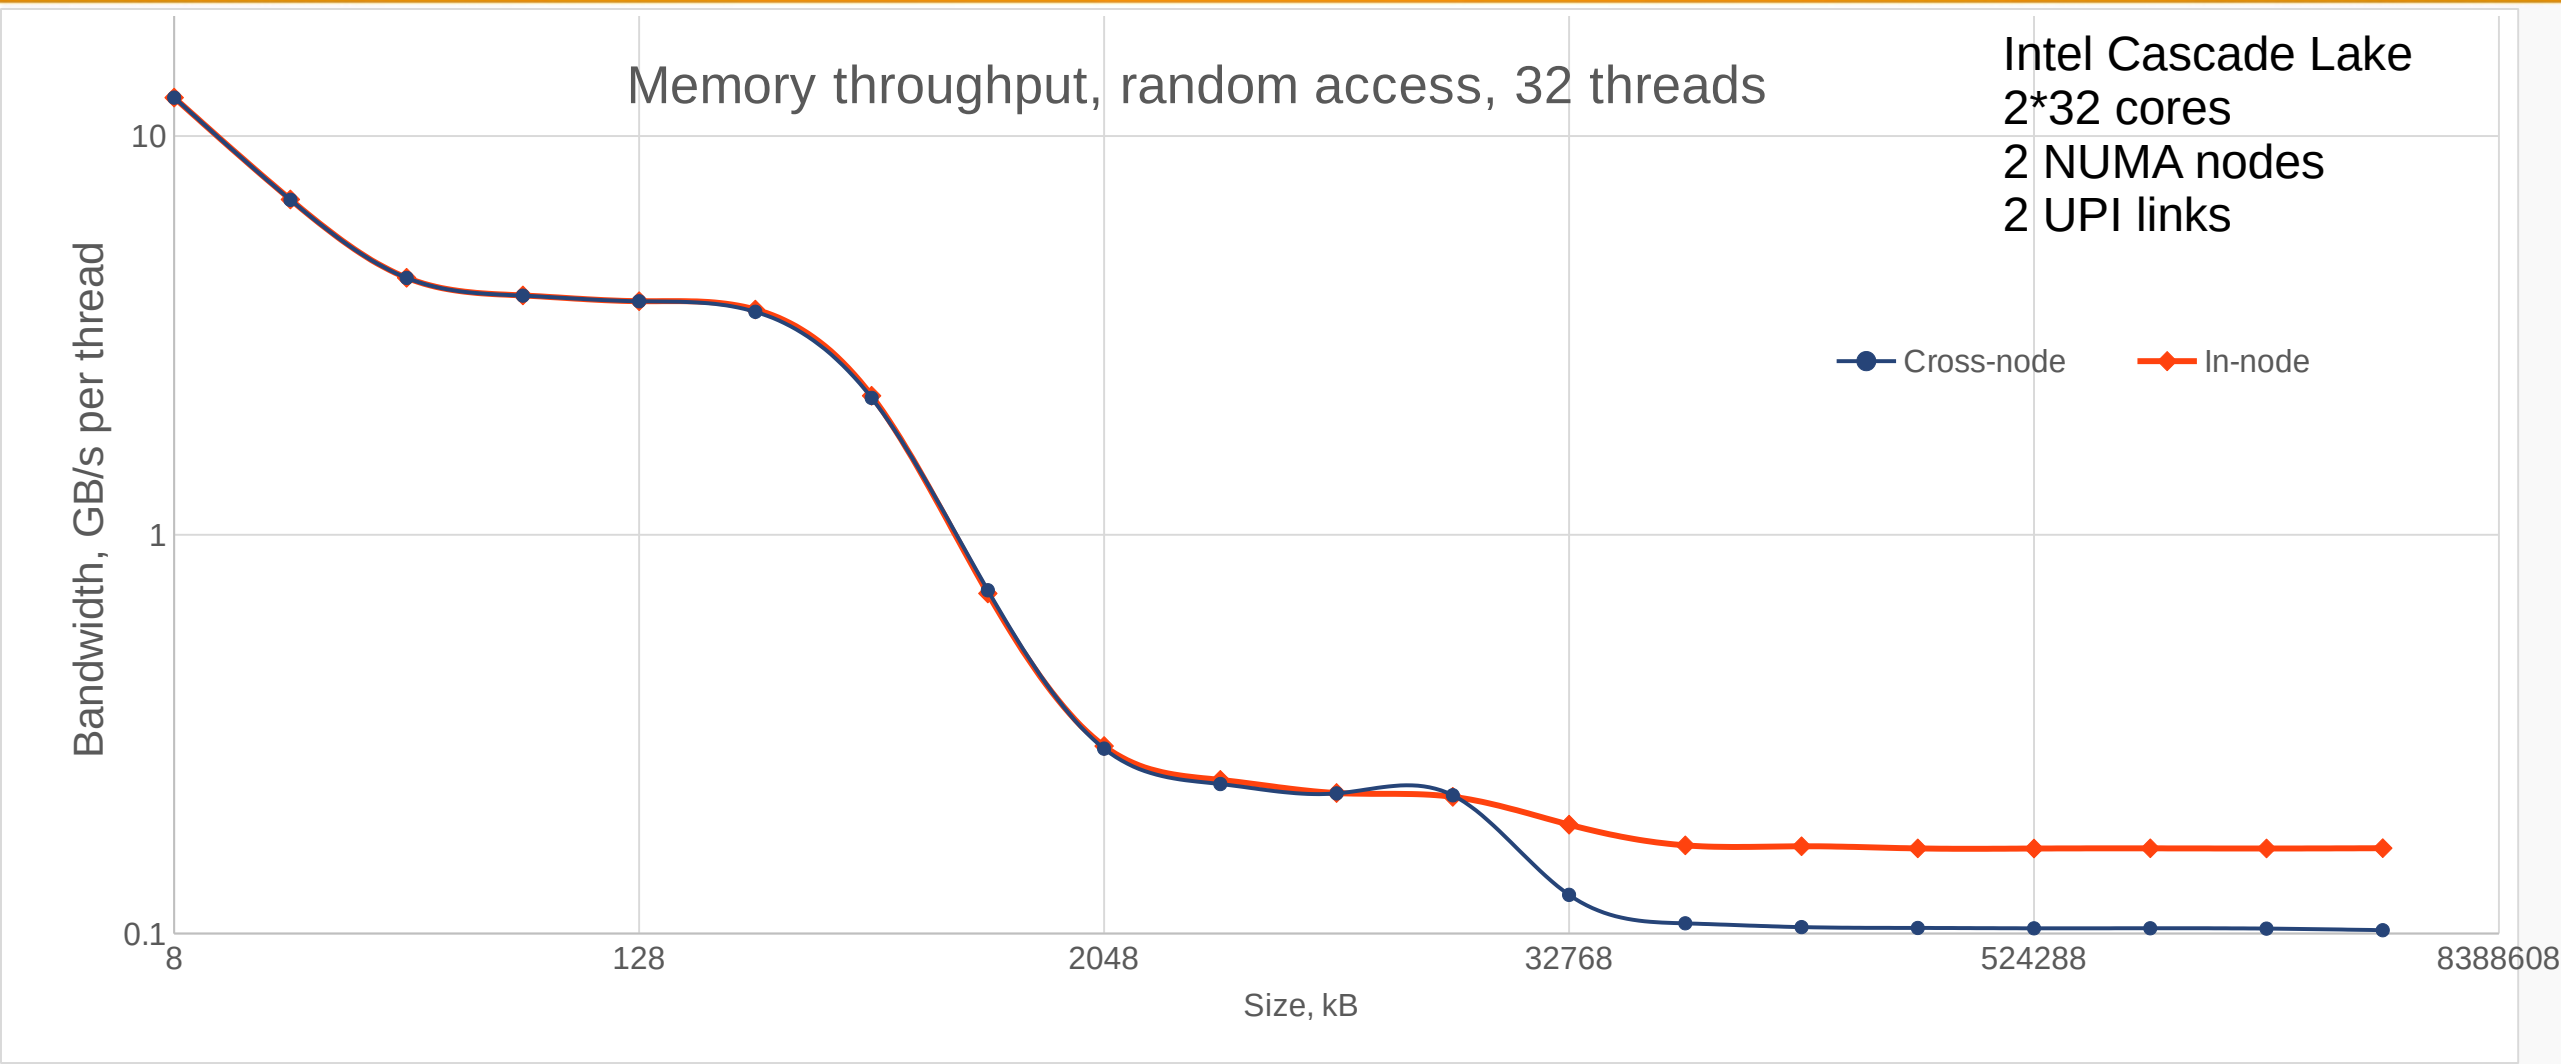
- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems



# Performance implications of NUMA

- Accessing local memory is different from accessing non-local memory
- To access memory from another node, that node's CPU must be involved
- Peak bandwidth is reduced by 40% to 60% for cross-node access
  - This is remarkably universal for x86 hardware from 2014 to 2022
  - Some systems show degradation for even 1 thread, some do not
  - Newer systems have higher bandwidth but larger penalty
- Cache bandwidth is not affected (without data sharing)
- Memory bandwidth is measured by sequential data transfers

# Memory Throughput in (newer) NUMA systems



# Performance implications of NUMA

- Accessing local memory is different from accessing non-local memory
- To access memory from another node, that node's CPU must be involved
- Peak bandwidth is reduced by 40% to 60% for cross-node access
  - This is remarkably universal for x86 hardware from 2014 to 2022
  - Some systems show degradation for even 1 thread, some do not
  - Newer systems have higher bandwidth but larger penalty
- Cache bandwidth is not affected (without data sharing)
- Memory bandwidth is measured by sequential data transfers
- Random access throughput is much smaller but cross-node penalty is similar
- All measurements are for independent memory access by each thread
  - Each thread allocates and writes its own memory

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is different from accessing non-local memory
- ~50% bandwidth penalty for remote node access (less if fits in cache)

# (Real) Performance implications of NUMA

- 50% memory throughput reduction is not too bad
  - May be observable in memory-bound programs
- All measurements are for independent memory access by each thread
  - What if all threads access the same data?


# (Real) Performance implications of NUMA

- 50% memory throughput reduction is not too bad
  - May be observable in memory-bound programs
- All measurements are for independent memory access by each thread
  - What if all threads access the same data? (for different values of “same”)
- Accessing the same address (true sharing) requires locks or atomic access
  - This is slow even between CPU cores within the same node
- Accessing different addresses within the same cache line (false sharing)
  - All the pain of sharing without the benefits (don't do it)
- Accessing different cache lines
  - This is usually as fast as any other memory accesses




# (Real) Performance implications of NUMA

- Test: all data resides on memory node 0
- 64 threads run on 64 cores (1 thread per core, all nodes)
- No cache lines shared between threads
- Thread data is interleaved and accessed sequentially (with stride)



# (Practical) Memory Bandwidth in NUMA systems



# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thead memory and cache locality are more important with NUMA
- Concurrent processing of data in close proximity may be slower


# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# (Real) Performance implications of NUMA

- Test: all data resides on memory node 0
- 64 threads run on 64 cores (1 thread per core, all nodes)
- No cache lines shared between threads
- Thread data is interleaved and accessed sequentially (with stride)
- Results depend strongly on timing of concurrent accesses

# (Practical) Memory Bandwidth in NUMA systems



# (Real) Performance implications of NUMA


- Test: all data resides on memory node 0
- 64 threads run on 64 cores (1 thread per core, all nodes)
- No cache lines shared between threads
- Thread data is interleaved and accessed sequentially (with stride)
- Results depend strongly on timing of concurrent accesses
- Running all 64 threads on the same NUMA node (only 32 cores) may be faster than using all 64 cores!




# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# Performance of a memory-bound program



# Observed performance without NUMA control



# If we do not control NUMA, who does?

- Unless explicitly overridden, the OS manages NUMA node loading
- CPU scheduler
  - Minimize thread migration between cores (usually)
  - Minimize thread migration between nodes (NUMA-aware scheduling)
  - Migrate threads to the node where the memory is (rarely)
- Memory allocation
  - Memory is usually allocated on the node running the thread (if available)
  - Memory may be migrated to another node if accessed from there (common)

# How to find out what you have? (Linux)

```
$ lscpu
CPU(s): 120
Thread(s) per core: 2
Core(s) per socket: 15
Socket(s): 4
NUMA node(s): 4
NUMA node0 CPU(s): 0-14,60-74
NUMA node1 CPU(s): 15-29,75-89
NUMA node2 CPU(s): 30-44,90-104
NUMA node3 CPU(s): 45-59,105-119
```



# How to find out what you have? (Linux)

```
$ numactl --hardware
```

```
available: 4 nodes (0-3)
```

```
node 0 cpus: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 60 61 62 63 64  
65 66 67 68 69 70 71 72 73 74
```

```
node 0 size: 386638 MB
```

```
...
```

```
node distances:
```

node	0	1	2	3
0:	10	21	21	21
1:	21	10	21	21
2:	21	21	10	21
3:	21	21	21	10



Memory latency (relative)

# How to find out what you have

- Intel memory latency checker
  - Not sure how relevant these numbers are...

```
$ mlc --latency_matrix -e -r
```

Measuring idle latencies for random access (in ns)...

Numa node

Numa node	0	1	2	3
0	130.7	187.9	193.4	200.8
1	190.2	129.0	192.2	194.0
2	196.0	194.4	130.1	192.9
3	203.4	192.1	193.9	129.4

# Simple NUMA controls (Linux)

`numactl --cpunodebind=N program args...`

- Run a program with all threads executing on node N (any CPU)

`numactl --membind=M program args...`

- Run a program with all memory allocations bound to node M
  - No memory transfer between nodes
  - Allocation will fail if the node is out of memory

# Debugging NUMA problems – step 0

- Run the program with NUMA restrictions and compare the results

```
$ numactl --cpunodebind=0 --membind=0 ./membench
```

```
64 threads      4194304 kB    2.160650 GB/s
```

```
$ numactl --cpunodebind=0 --membind=1 ./membench
```

```
64 threads      4194304 kB    0.670354 GB/s
```

- This program is sensitive to NUMA allocations, need to investigate...
  - Use reason – 1 node has half or quarter of all memory and CPUs
  - Running 64 threads on 32 cores (best NUMA access) may be slower than running 64 threads on 64 cores with NUMA penalties

# Debugging NUMA problems – step 0

- Run the program with/without NUMA restrictions and compare the results

```
$ ./parallel_foreach
```

```
32 threads: 4.74 million elements/second
```

```
$ numactl --cpunodebind=0 --membind=0 ./parallel_foreach
```

```
32 threads: 7.92 million elements/second
```

- This program is sensitive to NUMA allocations, need to investigate...
  - Hardware: 32 physical cores in 2 NUMA nodes (SMT – 2 logical cores per physical)



# Debugging NUMA problems – step 0

- Run the program with/without NUMA restrictions and compare the results

```
$ ./parallel_foreach
```

```
32 threads: 4.74 million elements/second
```

```
$ numactl --cpunodebind=0 --membind=0 ./parallel_foreach
```

```
32 threads: 7.92 million elements/second
```

```
$ numactl --cpunodebind=0 --membind=1 ./parallel_foreach
```

```
32 threads: 7.30 million elements/second
```

- This program is sensitive to NUMA allocations, need to investigate...
  - Hardware: 32 physical cores in 2 NUMA nodes (SMT – 2 logical cores per physical)

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- Simple benchmarks with different node binding often point to potential problems but usually not reveal their origin


# Debugging NUMA problems – step 1

- Profilers can help debug NUMA-related issues
- Indirect analysis: compare profiles for runs with different NUMA bindings
- Direct analysis: profile NUMA-related hardware counters
  - Intel VTune, Linux perf, PCM, etc
- Limited direct measurements
  - Can report QPI/UPI bandwidth per link
  - Large traffic between CPU sockets may indicate NUMA problems
- Profiling measurements are often misleading
- Linux kernel reports aggregate NUMA stats
  - Rarely accurate enough to be useful

# Practical debugging of NUMA problems

- If the program (or a part of it) runs faster when restricted to subset of nodes, you likely have a NUMA-related performance issue
  - Good breakdown of execution time is essential – instrument your code
- Compare profiles with and without NUMA restrictions or with favorable vs unfavorable NUMA restrictions
- Simplify the problem – benchmark suspected code fragments
  - Sometimes hardware profiles help (sometimes they mislead)
  - Often comparing results of numactl runs is your best approach
- Compare results on different hardware
  - Newer systems tend to expose NUMA-related problems

# Performance of a memory-bound program





# Comparing VTune profiles

	Interleaved	CPU-bound to node 0	
Instructions per cycle	0.04	0.08	Low (bound is better, but we knew that)
Memory-bound	42%		Memory-bound (we know that)
L1, L2, L3, DRAM-bound	2% to 0.5%		No obvious differences
NUMA remote access	50%	0.5%	Direct metric!
UPI-bound	50%	0%	Related metrics
UPI bandwidth	70%	20%	

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling sometimes can identify the problem and the location in the code

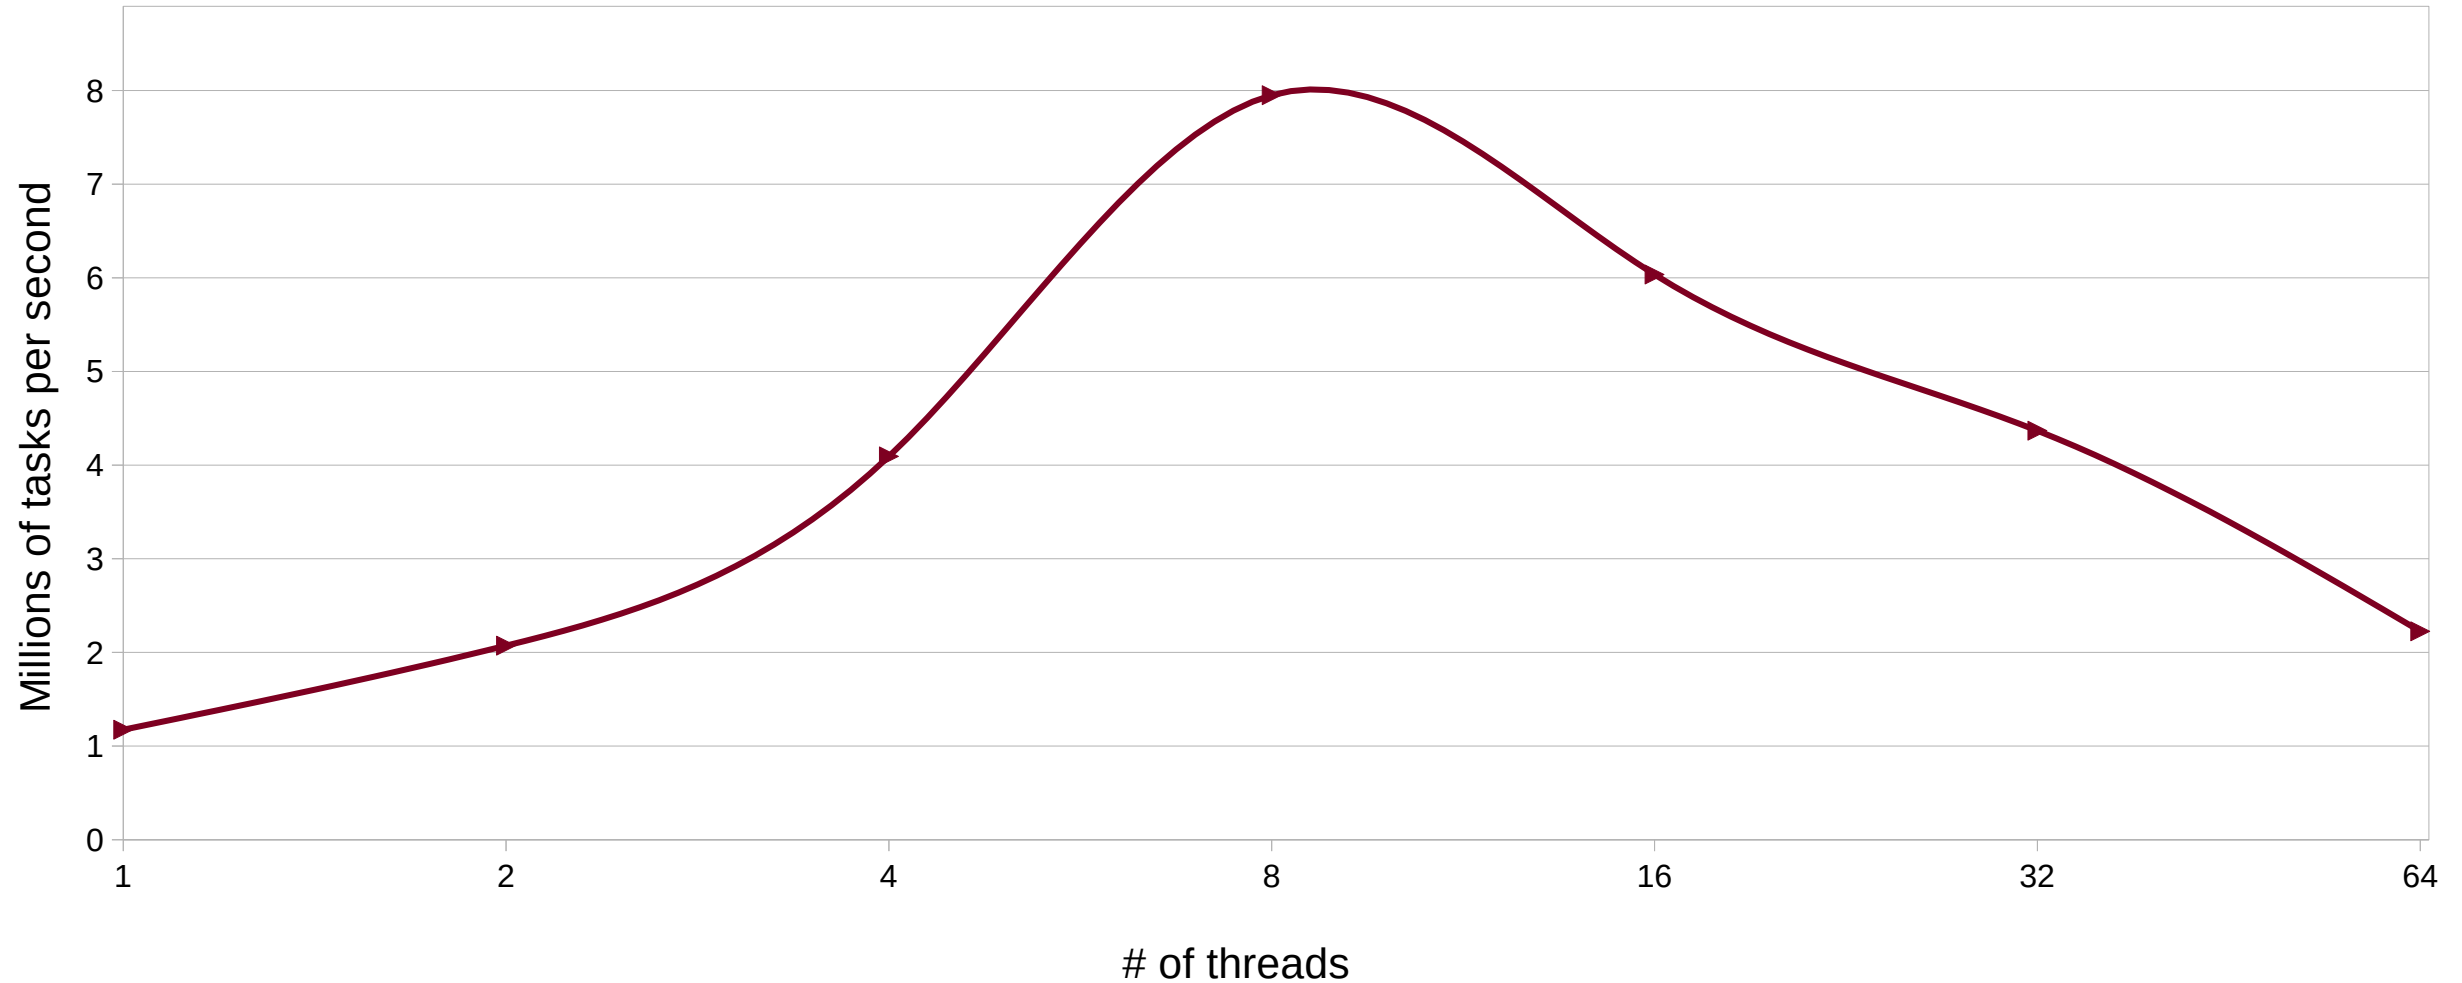
# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# History of this case...

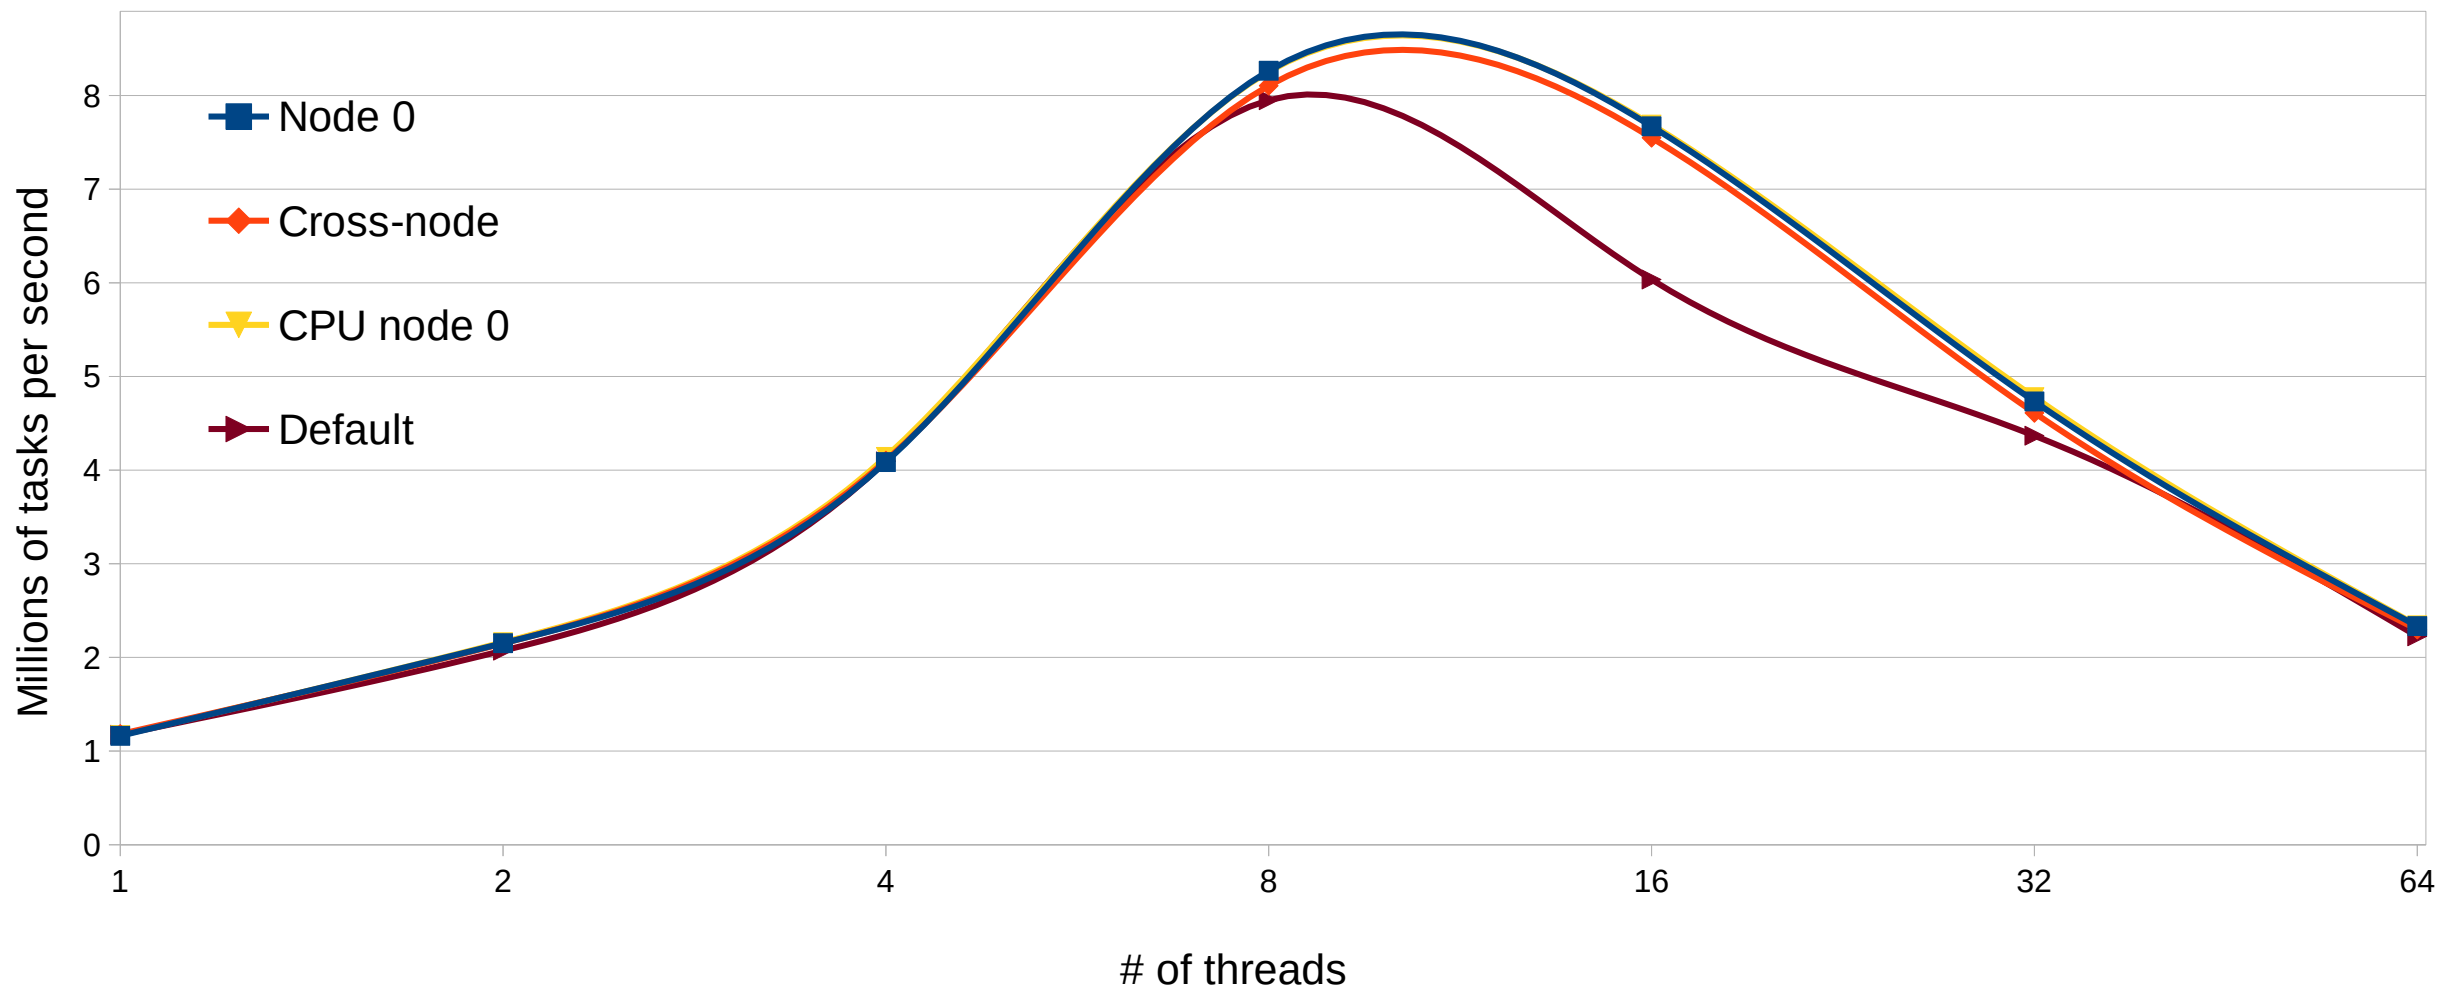
- Large multi-threaded program does not scale as much as expected
- Scaling is worse on newer processors
- Parts of the program run faster on one NUMA node than on two
- Problem step of the program is simplified and reduced to parallel foreach
- Per-element computations (parallel tasks) are compute-intense
  - Performance-critical code is not memory-bound
  - For simplification, memory use is further minimized

# Parallel foreach (no NUMA controls)





# Parallel foreach with numactl




# Performance analysis – first impression

- Restricting all threads to one node improves throughput
  - One node has 16 physical cores (2 SMT threads on each)
  - Running 32 threads on one node is faster than on both
- Restricting memory does not make any difference
  - Binding threads to node 0 and memory to node 1 is fast
  - Not binding threads is slower but test is unreliable, depends on the OS

# Performance analysis – first impression

- Restricting all threads to one node improves throughput
  - One node has 16 physical cores (2 SMT threads on each)
  - Running 32 threads on one node is faster than on both
- Restricting memory does not make any difference
  - Binding threads to node 0 and memory to node 1 is fast
  - Not binding threads is slower but test is unreliable, depends on the OS
- If you want something done right, do it yourself
  - Bind half of all threads to node 0 and half to node 1

# Parallel foreach



# Performance analysis – second impression

- Memory is irrelevant
  - Makes sense – the simplified program uses very little memory
- Forcing threads to run on both nodes equally is even worse than whatever the OS does
- Restricting all threads to one node makes the program run faster
  - Makes no sense – there aren't enough cores in one node for all threads
- Time to use the profiler



# Profiling parallel foreach

- Profile the worst (strangest?) case – 32 threads on one vs two nodes
- Try different hardware counters until enlightenment achieved
- VTune's hints may be helpful or misleading
- Profiling the problem case by itself is rarely useful
- Comparing two versions with systematic differences is often productive

# Profiling with perf – original code

- 617276143376 cycles:u  
21575458142 instructions:u # 0.03 insn per cycle  
541435483399 cycle\_activity.stalls\_mem\_any:u
- 1309864032 cache-references:u  
459547539 cache-misses:u # 35.084 % of all cache refs
- 281154229 LLC-loads:u  
85821152 LLC-load-misses:u # 30.52% of all LL-cache accesses
- 135375847 LLC-stores:u  
41651753 LLC-store-misses:u # 30.77% of all LL-cache accesses
- Also significant fraction of front-end stalls and other problems

# Profiling with perf – CPU bound to node 0

- 245669779655 cycles:u  
38546010583 instructions:u # 0.16 insn per cycle  
114854490177 cycle\_activity.stalls\_mem\_any:u
- 1536581119 cache-references:u  
54666032 cache-misses:u # 3.558 % of all cache refs
- 232443098 LLC-loads:u  
3786766 LLC-load-misses:u # 1.63% of all LL-cache accesses
- 238253412 LLC-stores:u  
3184831 LLC-store-misses:u # 1.33% of all LL-cache accesses
- Program uses very little memory, writes ~MB/s (vs ~10 GB/s bandwidth)

# Comparing VTune profiles

	Unrestricted	CPU-bound to node 0	
Instructions per cycle	0.18	0.24	Low (bound is better, but we knew that)
CPU utilization	20%	30%	Low; causes: memory stalls, branch misses, instruction starvation
Front-end stalls	52%	30%	High, causes: d-cache misses, i-cache misses
Memory-bound	34%	19%	High, mostly cache-bound (why?)
NUMA remote access	80%	20%	Direct metric, but of what?

- The program is very small and uses very little memory

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling can be misleading or boil down to “something somewhere is somehow gone wrong”



# What do we know? What makes sense?

- The benchmark is a parallel foreach (4M of code including C++ runtime)
- Each task is about 1us, uses 8B of memory
- On NUMA systems, scaling is poor when thread count exceeds number of CPUs in a single node
- Binding all threads to one NUMA node improved throughput slightly
  - Even if the node doesn't have enough CPUs
- Binding memory to one NUMA node makes no difference
- Profilers report large percentage of LLC misses
  - The entire data set fits into LLC
- Profilers report front-end problems usually seen in much larger code


# What do we know? What makes sense?

- The problem is worse in newer systems (UPI scales worse than QPI)
  - UPI has higher bandwidth
- Time to ask for help from the people who designed the processor
  - Latest version of UPI has even higher bandwidth (and runs even slower)
  - The program is latency-bound
- QPI cross-node latency: 120-140 ns (40% over in-node)  
UPI cross-node latency: 140-160 ns (80% over in-node)
- Newer CPUs have higher interconnect bandwidth and higher latency
  - Both relative and absolute latency

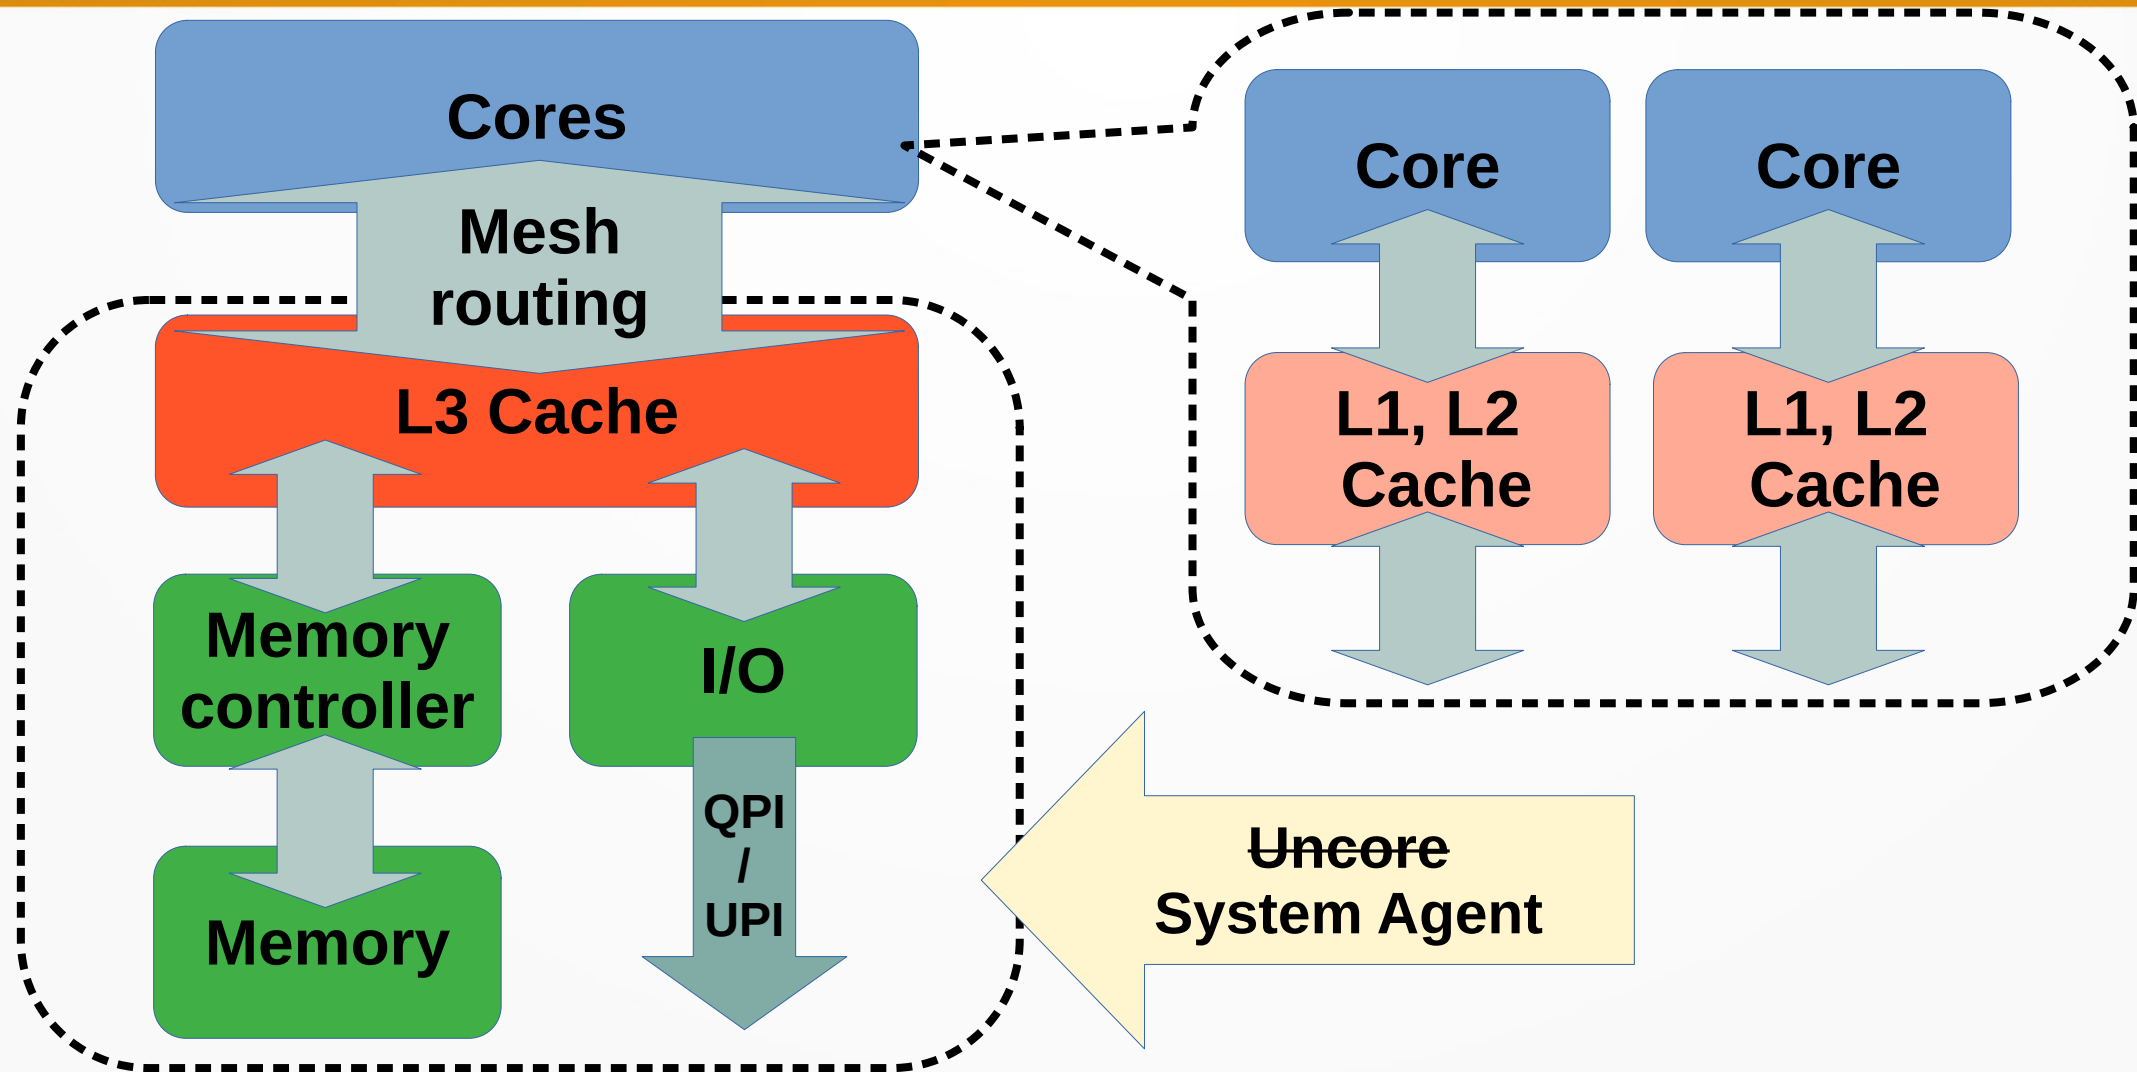
# Why (when) latency is important?

- Random access of small amounts of data
  - Not relevant for our problem
- Concurrent access of the same data – data sharing
  - Throughput is limited by the latency of atomic accesses
- Foreach is perfectly parallel, no data sharing at all, but...
- Any thread scheduler is going to have some shared data
  - Task queue (queue lock or atomics)
  - Task count or whatever you wait on
  - Other data, depending on the thread scheduler/pool implementation

# Multi-threading and latency

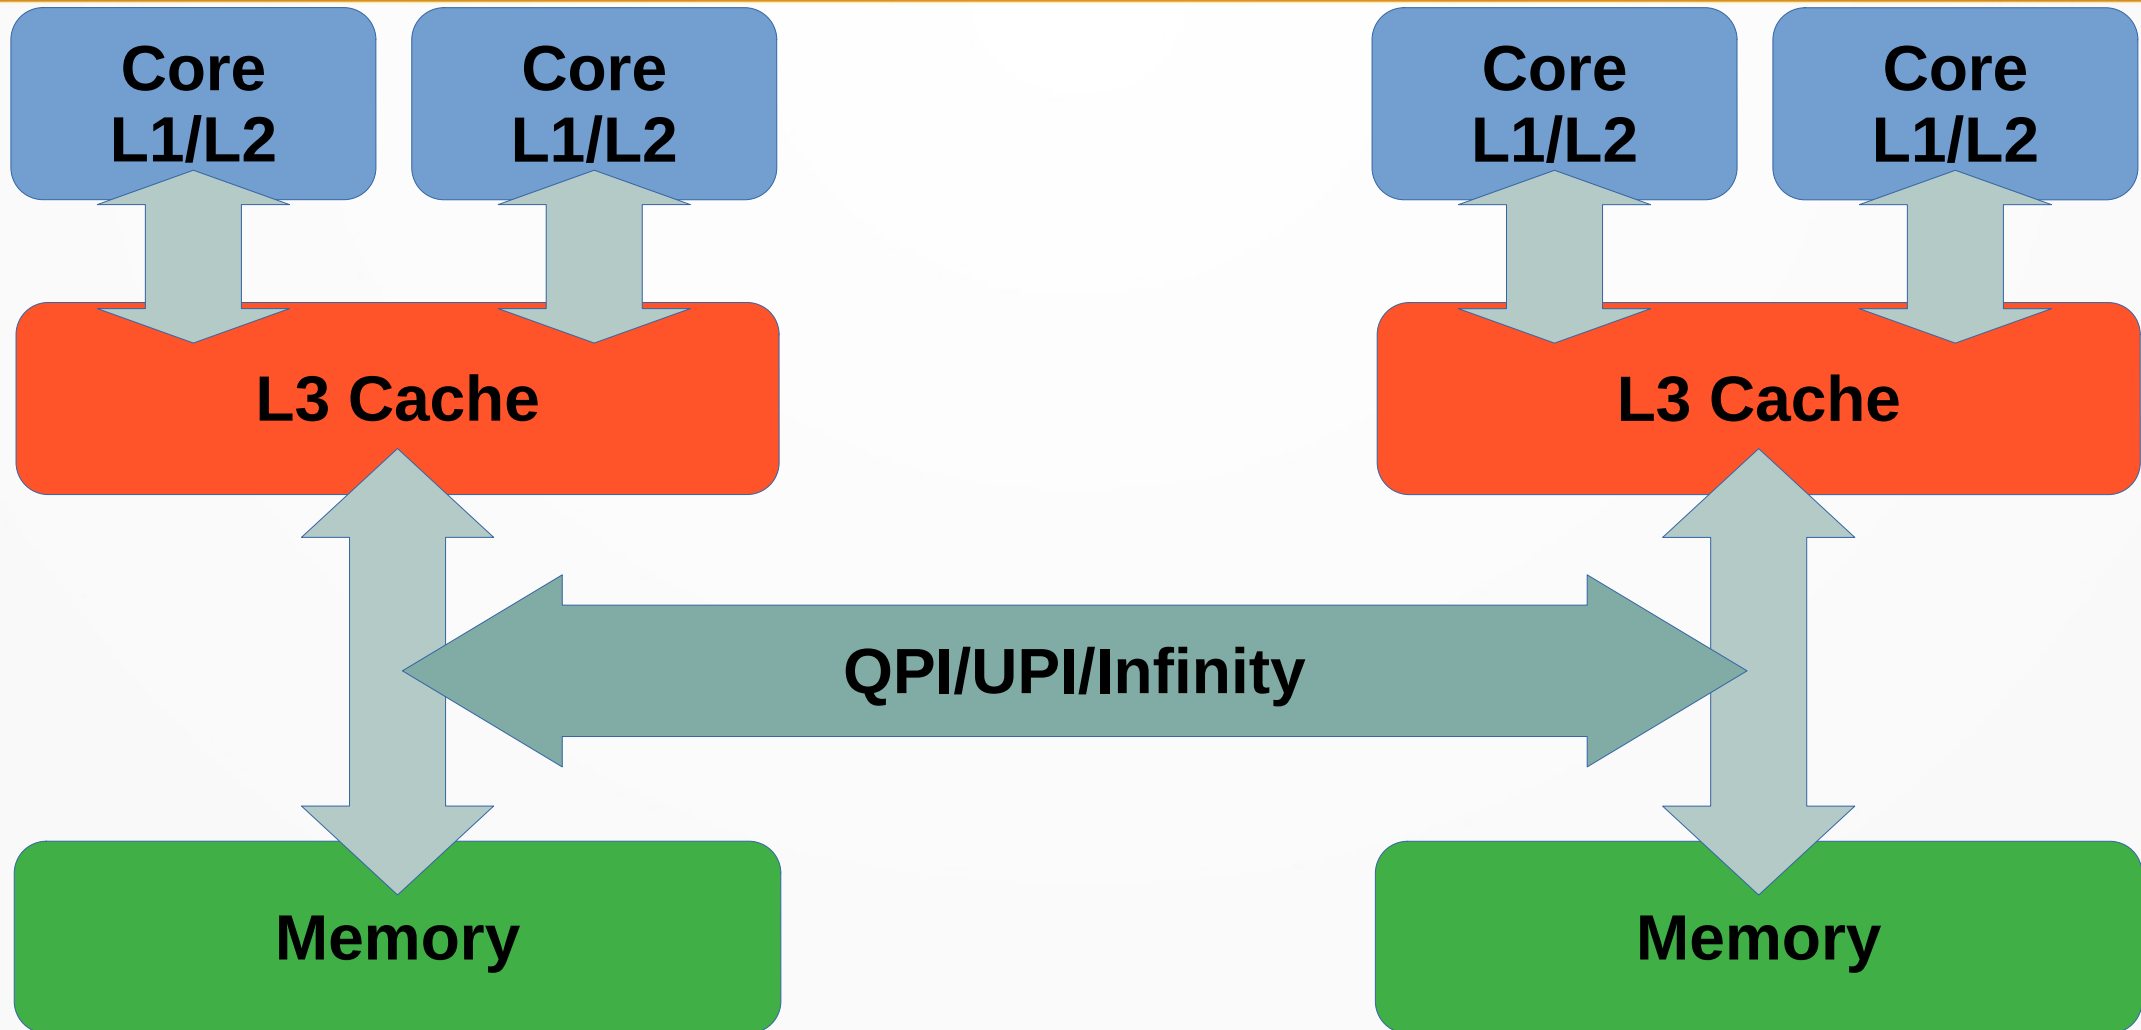


# Multi-threading and latency






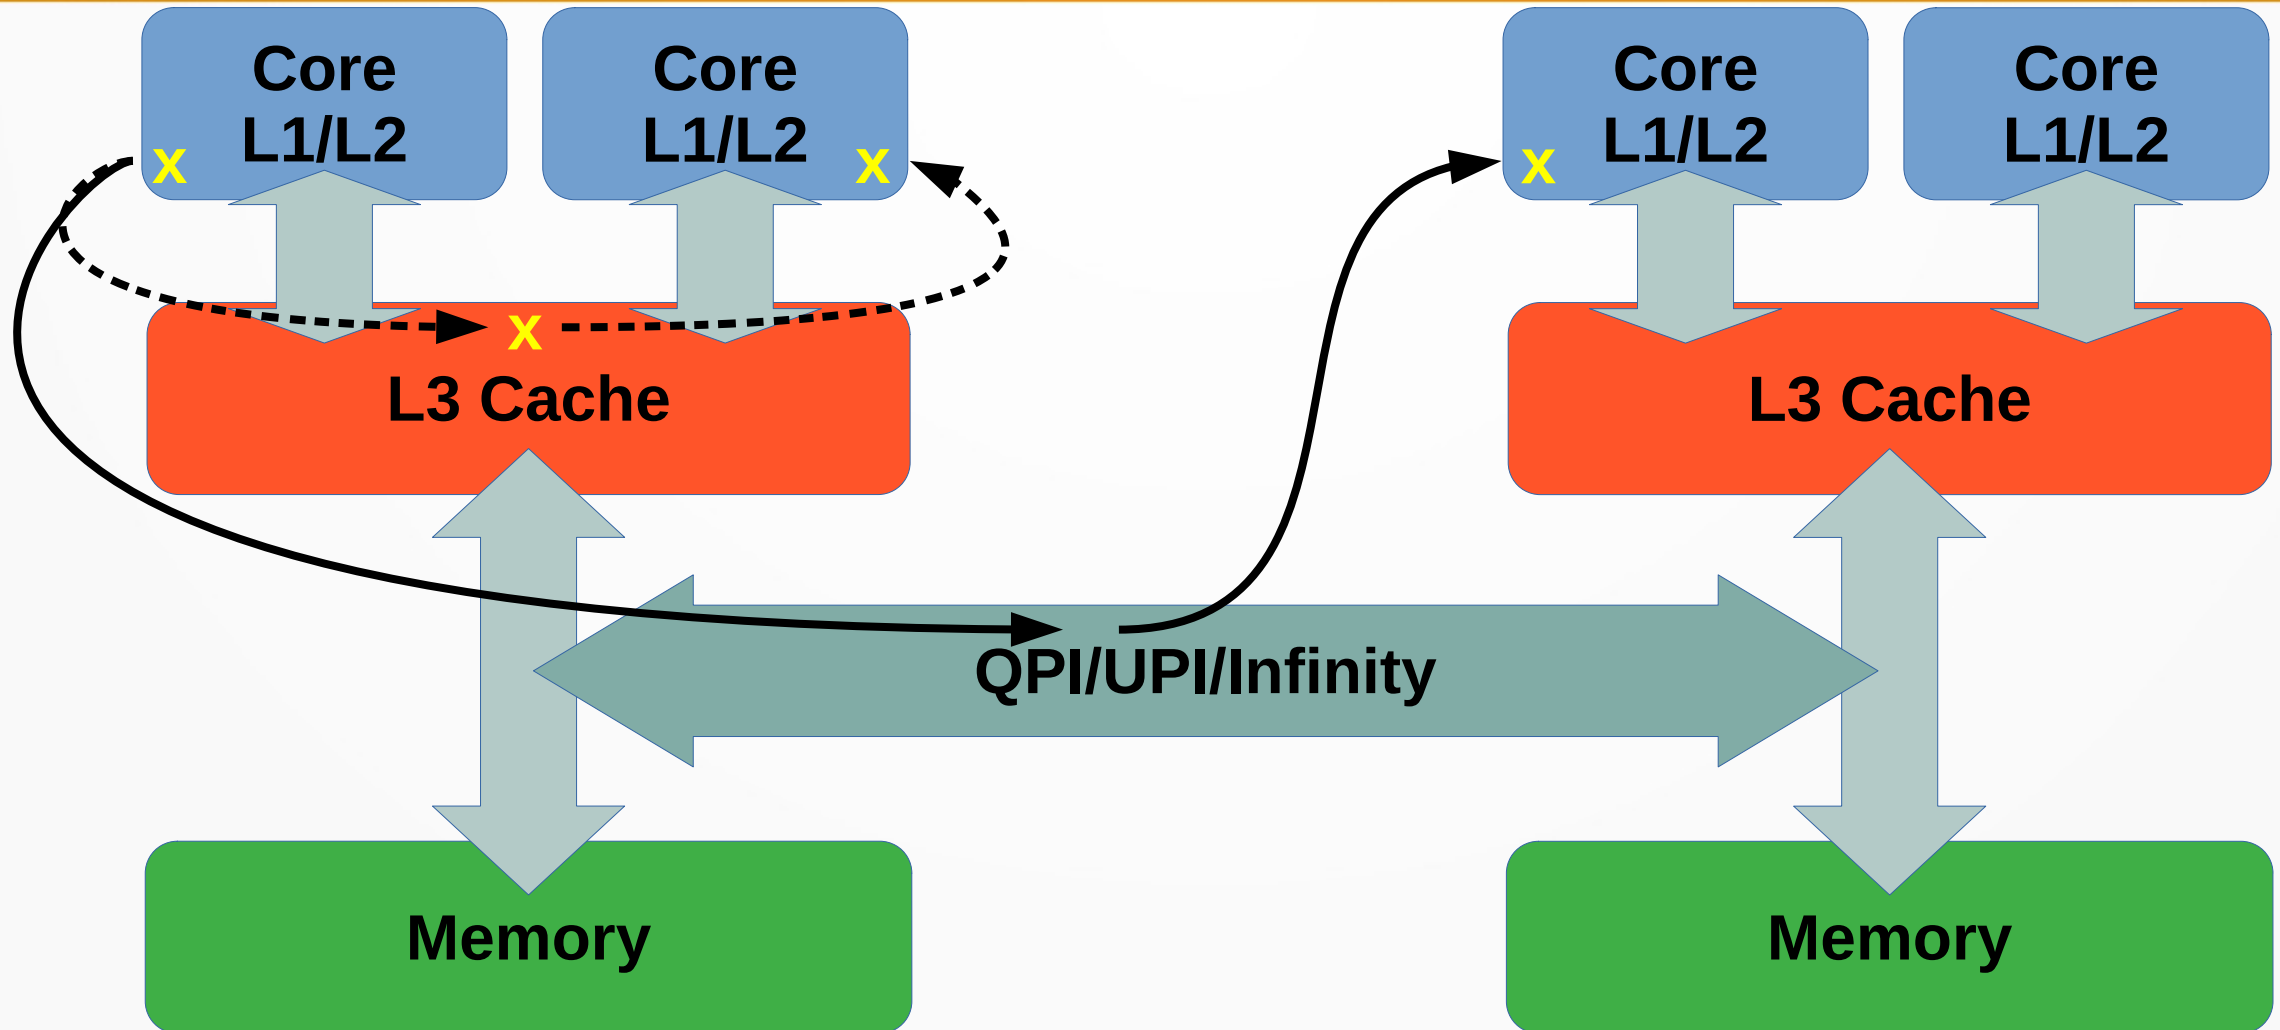
# Multi-threading and latency




# This is your data sharing



# This is your data sharing on NUMA



# Is NUMA performance hurt by data sharing?



# Is NUMA performance hurt by data sharing?

- Performance of a single atomic operation often shows some slowdown on NUMA systems
  - Especially for Read-Modify-Write operations
  - Usually not for atomic load or store (so RCU algorithms may have advantage)
  - Intel QPI systems may not show any slowdown at all
  - AMD Epyc systems show little to none NUMA penalty
- Most concurrent data structures and executors use shared variables and atomic operations
  - Locks too (usually atomic exchange)
- Maintaining consistent global shared state on NUMA systems is expensive

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- True data sharing can be significantly more expensive
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling is useful if you know what to look for



# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# NUMA control beyond numactl

- NUMA programming API allows very fine control of thread and memory interactions with NUMA hardware

```
#include <numa.h> // Link with -lnuma
```

```
#include <numaif.h>
```

- Restrict the calling thread to a subset of NUMA nodes or CPUs:

```
numa_run_on_node(), numa_sched_set_affinity()
```

also pthread\_setaffinity\_np() from pthreads library

- Restrict memory allocations by the calling thread to specific NUMA nodes:

```
numa_set_preferred(), set_mempolicy()
```


- Functions for direct allocations on the specified node, moving memory between nodes, querying current NUMA policies and hardware capabilities
- Caution: cpu-to-node mapping is SLOOOW (cache cpu maps and masks)

# When to use NUMA controls?

- In many applications, never (it just works)
- Areas of concern:
  - High-throughput or low-latency concurrent data structures
  - Thread pools, executors, schedulers, etc
  - Memory-bound programs
- Try to bind the program to one NUMA node (or bind CPU and memory to different nodes) to see if performance is affected
- If you must use NUMA API, encapsulate it (NUMA pool, NUMA counter)
  - Exception: NUMA-friendly memory management may impact the way you do concurrency in your algorithm


# Memory-bound programs

- If each thread mostly uses its own memory, you can bind threads to NUMA nodes (but be mindful of load balancing)
  - Verify that it is needed – OS scheduler usually does good enough job




# Memory-bound programs

- Memory is usually allocated on the calling thread's node by default
  - But if the parent thread allocates memory, you may want to explicitly move it
  - Recent Linux kernels will eventually do it for you (with some caveats)



# Memory-bound programs

- If multiple threads update the same data structure (but not the same memory locations) the working sets of these threads must be separated
  - Read-only data is OK
- Algorithm changes may be required to increase per-thread work size
  - Separated by how much? One page (4K) may be enough





# Memory-bound programs


- If each thread mostly uses its own memory, you can bind threads to NUMA nodes (but be mindful of load balancing)
  - Verify that it is needed – OS scheduler usually does good enough job
  - Memory is usually allocated on the same node by default
  - But if the parent thread allocates memory, you may want to explicitly move it
  - Recent Linux kernels will eventually do it for you (with some caveats...)
- If multiple threads update the same data structure (but not the same memory locations) the working sets of these threads must be separated
  - Read-only data is OK
- Algorithm changes may be required to increase per-thread work size
  - Separated by how much? One page (4K) may be enough (but there are complications...)

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- True data sharing can be significantly more expensive
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling is useful if you know what to look for
- Good practices for memory-bound programs are more important
- Consider binding threads and/or memory if OS defaults are insufficient

# Latency-bound programs


- If cross-node traffic caused by some shared data is significant, the implementation has to be redesigned to use per-node state



- But we still need a global state...

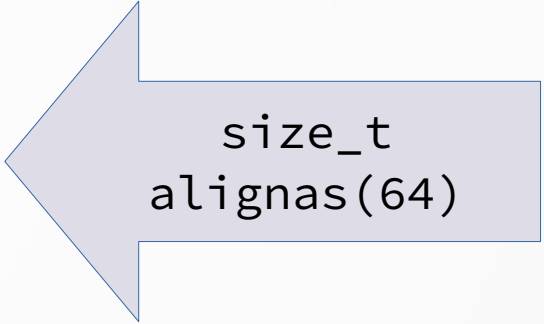
# NUMA-aware counter

- If cross-node traffic caused by some shared data is significant, the implementation has to be redesigned to use per-node state
- Synchronization is likely to be more complex

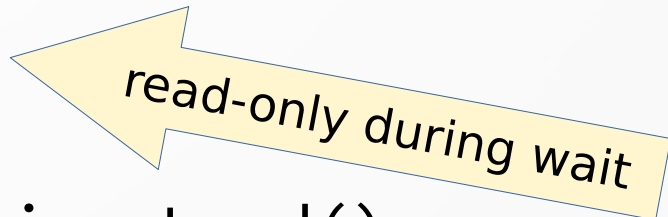


# NUMA-aware counter

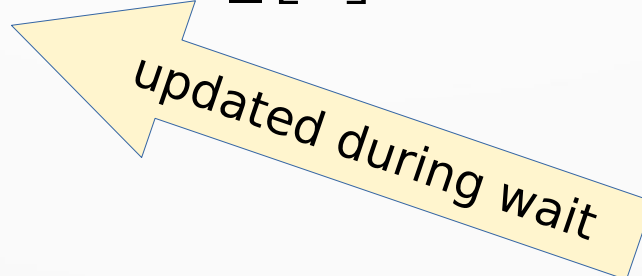
```
class NUMA_task_count {  
    static constexpr size_t max_node_count = 64;  
    task_count_t done_task_counts_[max_node_count];  
    std::atomic<size_t> task_count_;  
public:  
    size_t Count() const { // This is what we wait on  
        size_t res = task_count_.Acquire_Load();  
        for (size_t i = 0; i != num_nodes; ++i)  
            res -= done_task_counts_[i].value.Acquire_Load();  
        return res;  
    }  
}; // NUMA_task_count
```



size\_t  
alignas(64)




read-only during wait



updated during wait

# NUMA-aware thread pool


- Per-node pools, cross-node work stealing (only when one node is empty)
- Task submission from the main thread is the bottleneck






# NUMA-aware thread pool

- Per-node pools, cross-node work stealing (only when one node is empty)
- Task submission from the main thread is the bottleneck
- Can be solved with a two-stage task submitter




# NUMA-aware thread pool

- Per-node pools, cross-node work stealing (only when one node is empty)
- Reduce cross-node data contention – fewer interacting threads or less frequent interaction



# NUMA-aware thread pool



# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- True data sharing can be significantly more expensive
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling is useful if you know what to look for
- Use good practices for memory access and node binding only if needed
- [Some] Concurrent data structures need to be redesigned for NUMA

# We Will Learn:

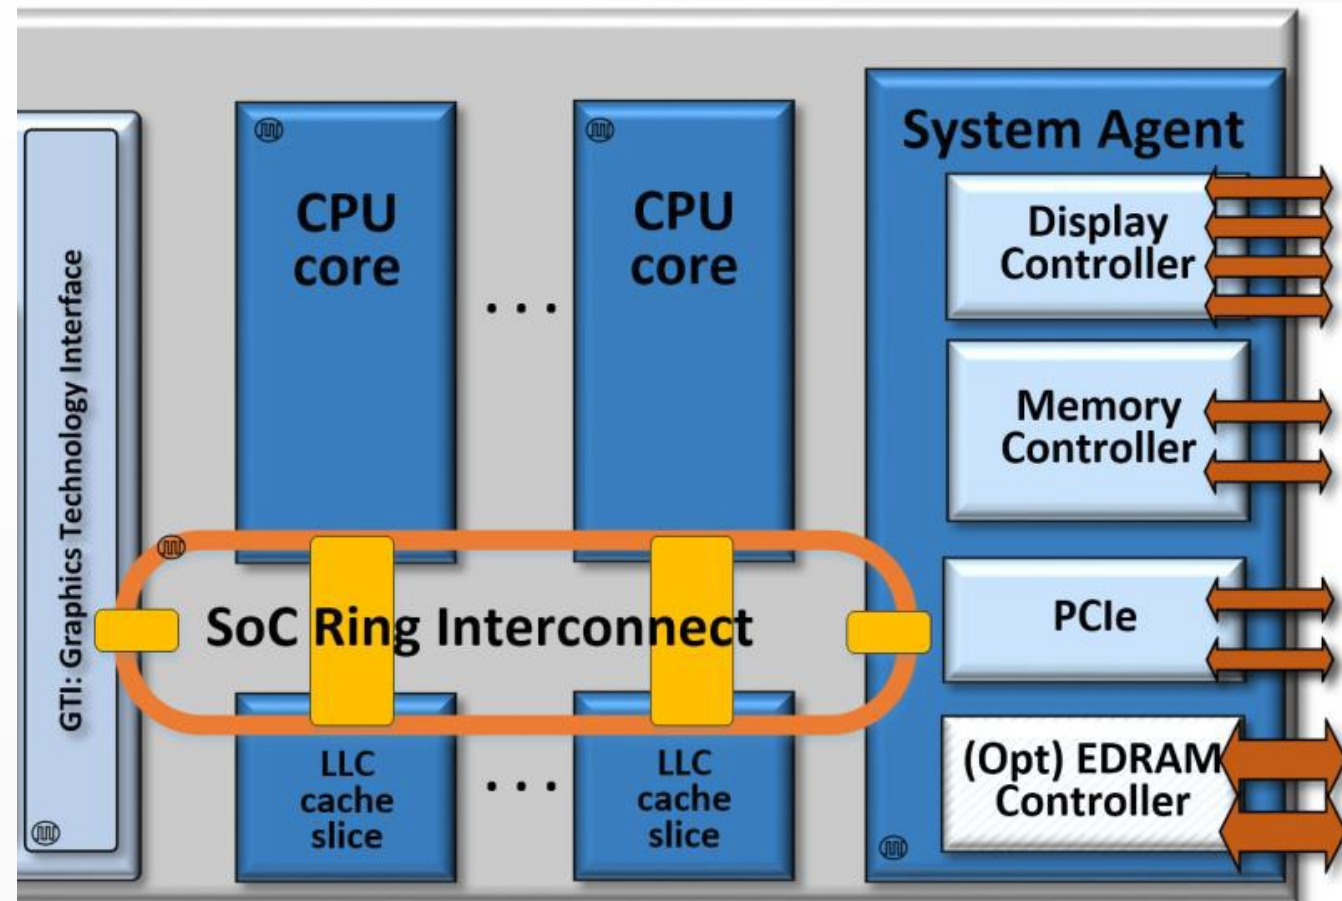
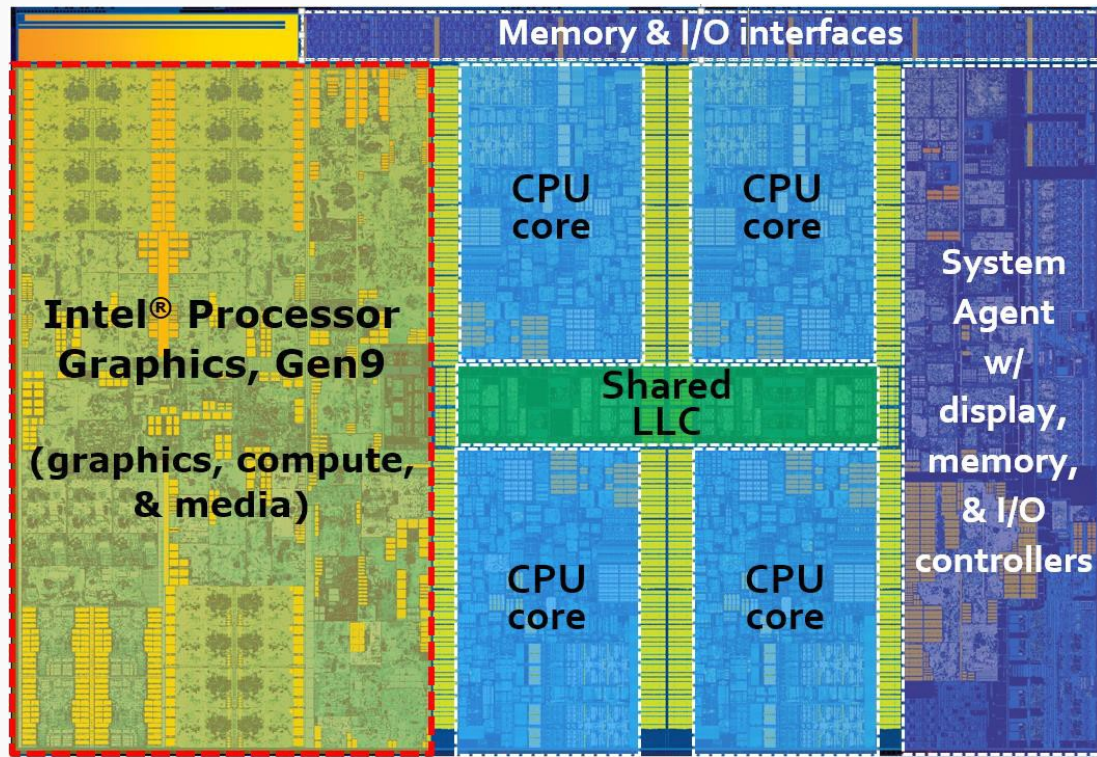
- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems

# NUMA concerns for distributed applications



# NUMA concerns for distributed applications

- Distributed systems often send a lot of traffic across the network



# NUMA concerns for distributed applications

- Distributed systems often send a lot of traffic across the network
- Which NUMA node is the network controller connected to?

```
# /sbin/lspci
```

```
04:00.0 Ethernet controller: Intel Corporation 82599ES
```

```
# cat /sys/bus/pci/devices/0000\:04\:00.0/numa_node
```

```
0 
```

- Does it matter? Sometimes.
- Case study: distributed application (1000s of remote CPUs)
  - Use numactl to bind to node 0 vs node 1 – 10% run time difference
  - Built-in throughput tester: 20% difference between nodes 0 and 1

# NUMA concerns for I/O-bound applications

- High-speed PCI-E devices are attached to a specific NUMA node
- NUMA may be important if the device is very fast
  - Simple test: bind the program to node 0 or node 1 and compare
    - Comparing with unbound run mostly tests the OS scheduler
- No case study – none of our programs are sufficiently I/O-bound



# NUMA concerns for GPUs and accelerators

- GPUs are a special case of I/O
  - GPU interfaces are usually faster than any other device
  - Data transfer is often the bottleneck of GPU acceleration
- CUDA Bandwidth test for Tesla V100: 10GB/s (node 0) vs 7GB/s (node 1)
- BUT...
- If GPU-CPU transfer is a concern, pinned memory should be used
- Pinned memory is on the GPU node
  - At least by default?
  - Transfer rate 13GB/s

# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- True data sharing can be significantly more expensive
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling is useful if you know what to look for
- Use good practices for memory access and node binding only if needed
- [Some] Concurrent data structures need to be redesigned for NUMA
- For I/O-bound programs, use the right node for the data

# We Will Learn:

- What is NUMA
- Why NUMA matters for programmers
- Performance implications of NUMA
  - Straightforward
  - Non-intuitive
  - Weird
- How to debug NUMA issues
- How to program for NUMA systems



# Real-life case

- Program runs slower on (faster) NUMA hardware than on old machines
- The slowdown is limited to certain places in the code
- For these functions, slowdown is large (2x to 10x)
- The code is not limited by bandwidth or latency but there is one common trait: they touch large amounts of memory
  - Like updating links in a list of large elements
- The slowdown is OS-dependent:
  - Not present on old versions of Linux (overall, newer OS → faster runs)

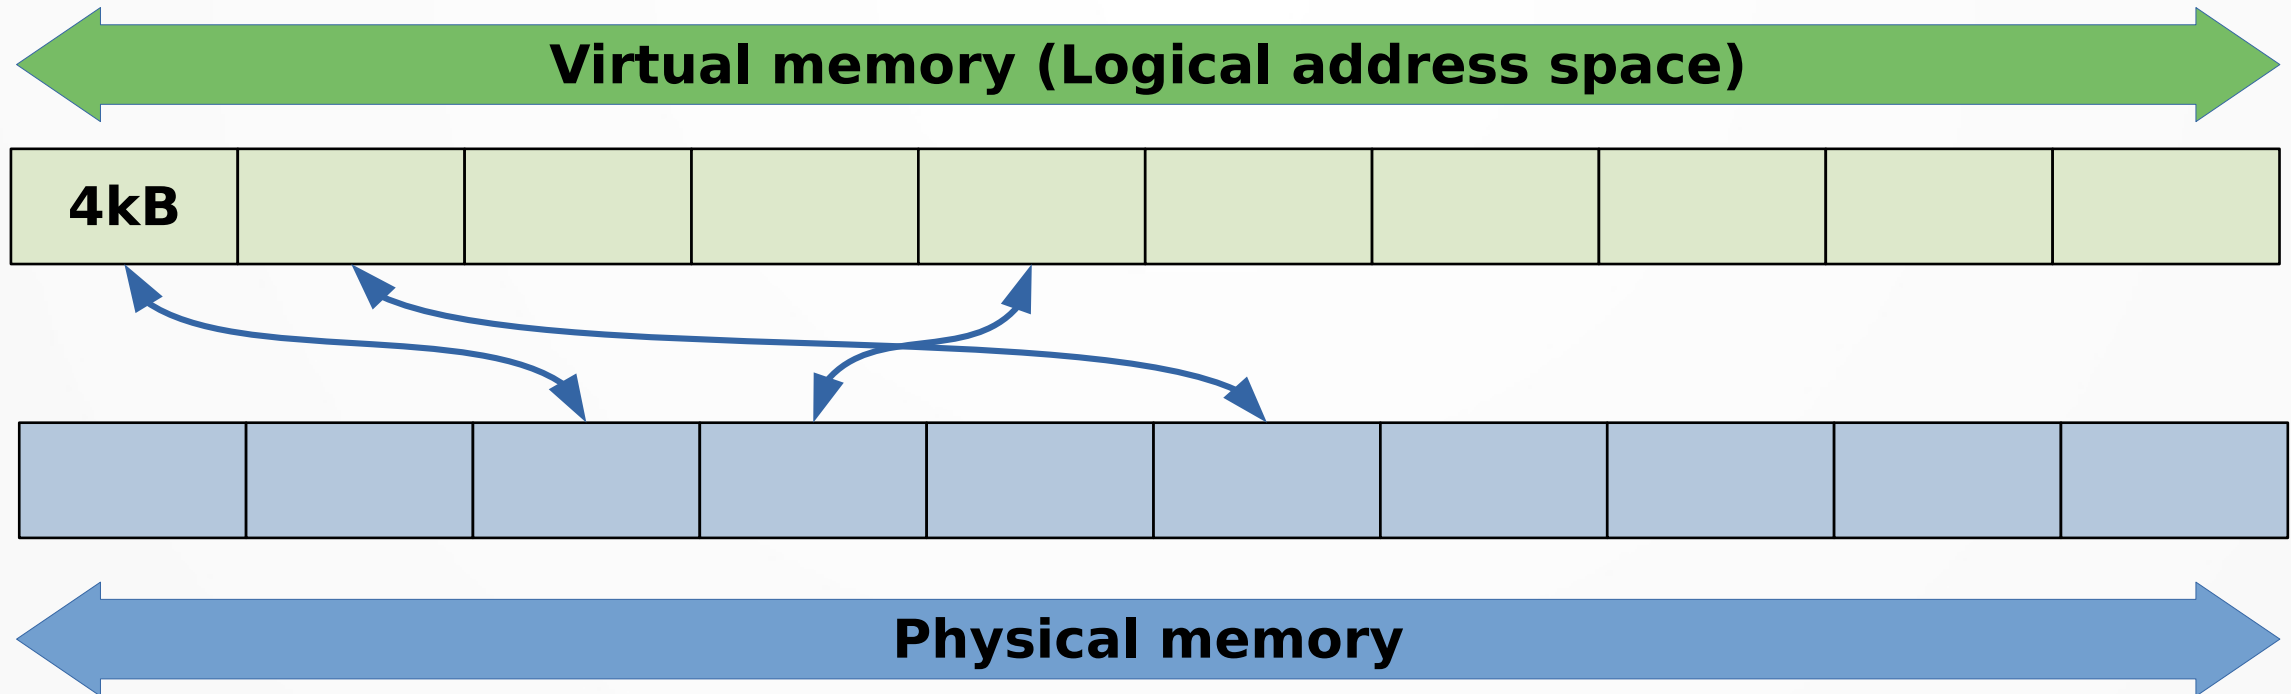
# Debugging of the slowdown

- NUMA-related: the execution steps in question run faster when bound to one NUMA node
- Not easy to reproduce: specific memory access patterns are essential, hard to reproduce artificially
  - Plan B: modify the real program to run the problem step repeatedly or on larger data volume
- Profile with hardware counts: cache misses typical for semi-random memory access but no worse than elsewhere
- Built-in time reporting: real time is longer than CPU time
  - Even with multi-threading
  - There is no I/O or network traffic
  - Detailed time reporting: mostly system CPU time, not user time

# Debugging of the slowdown

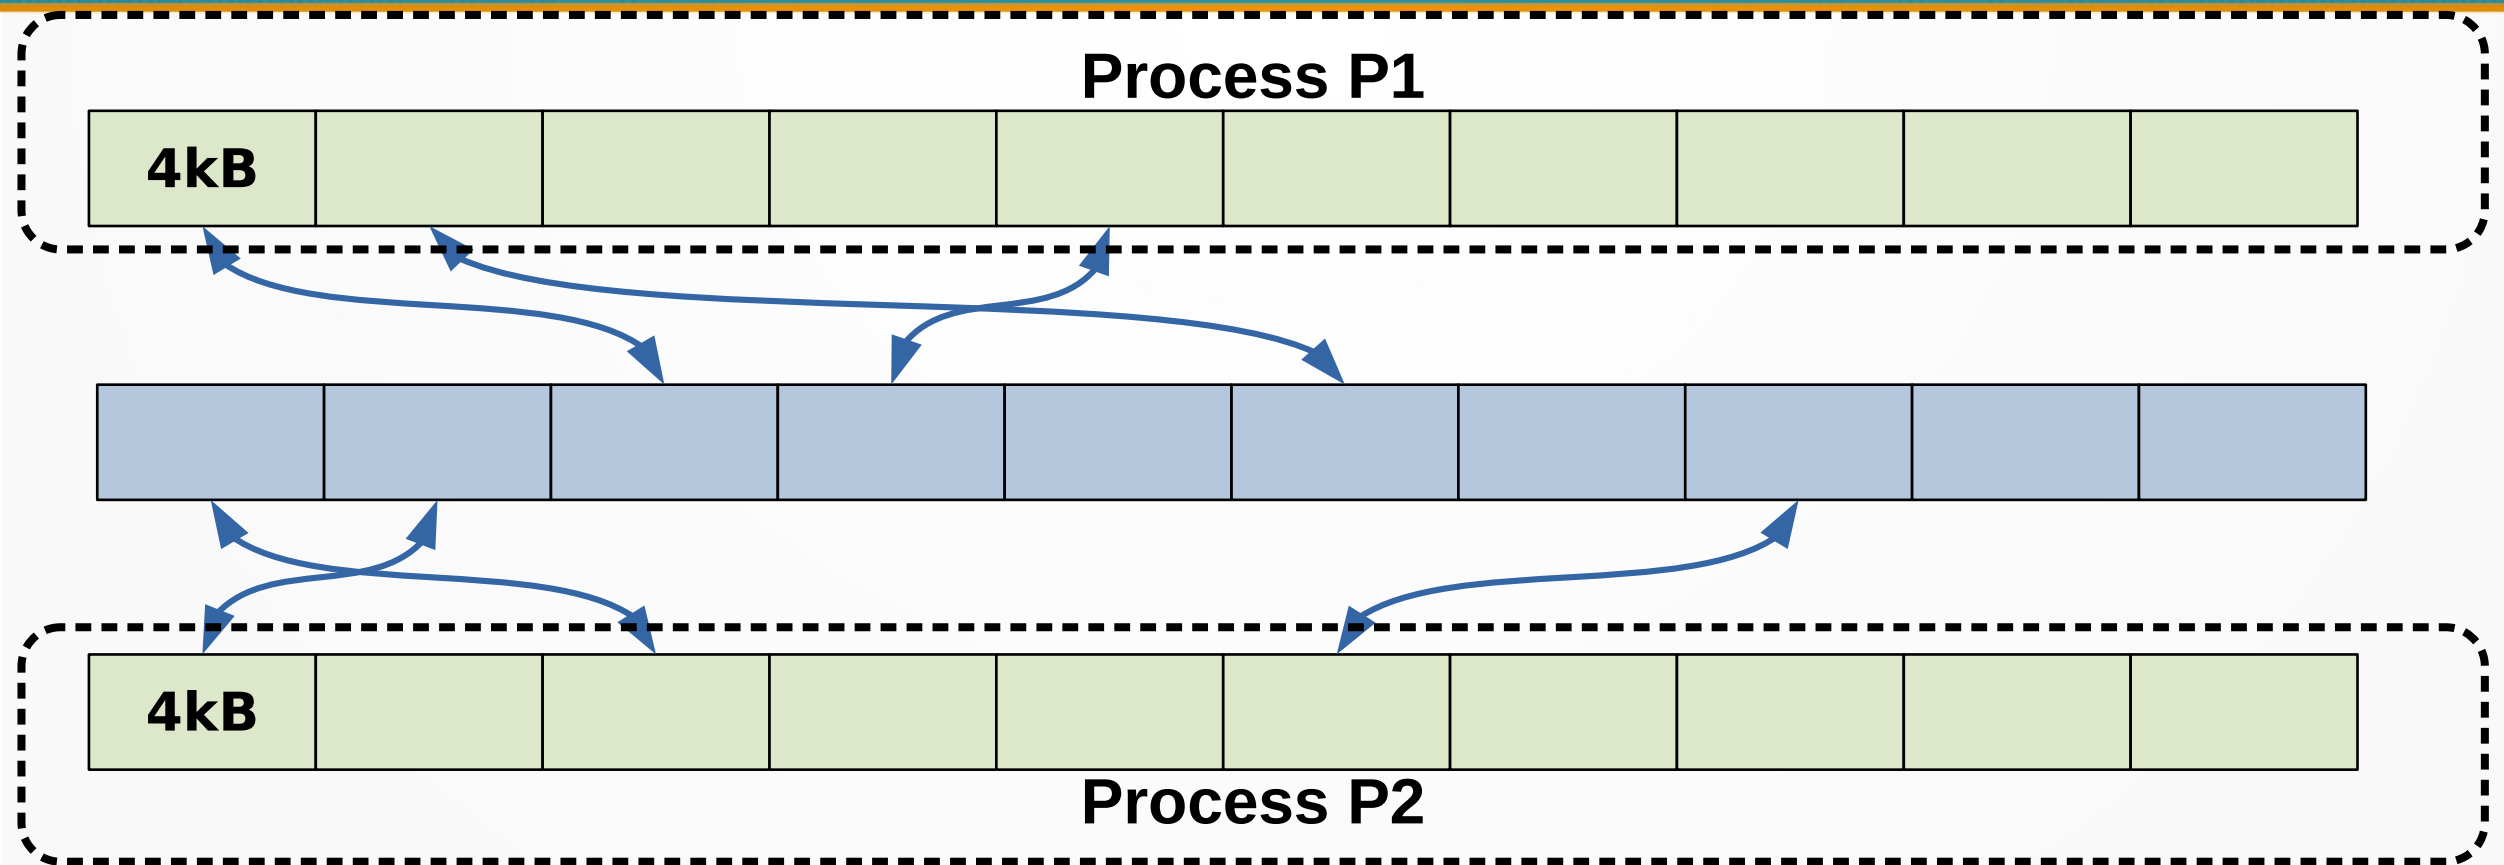
- Debug setup: modify the real program to run the problem step repeatedly or on larger data volume
- Detailed time reporting: mostly system CPU time, not user time
- Profiling shows that the time is spent in the Linux kernel
- Profiling with kernel symbols shows that the time is spent in two sections:
  - NUMA management (memory migration, etc)
  - TLB updates
- What is TLB?

# Virtual memory management



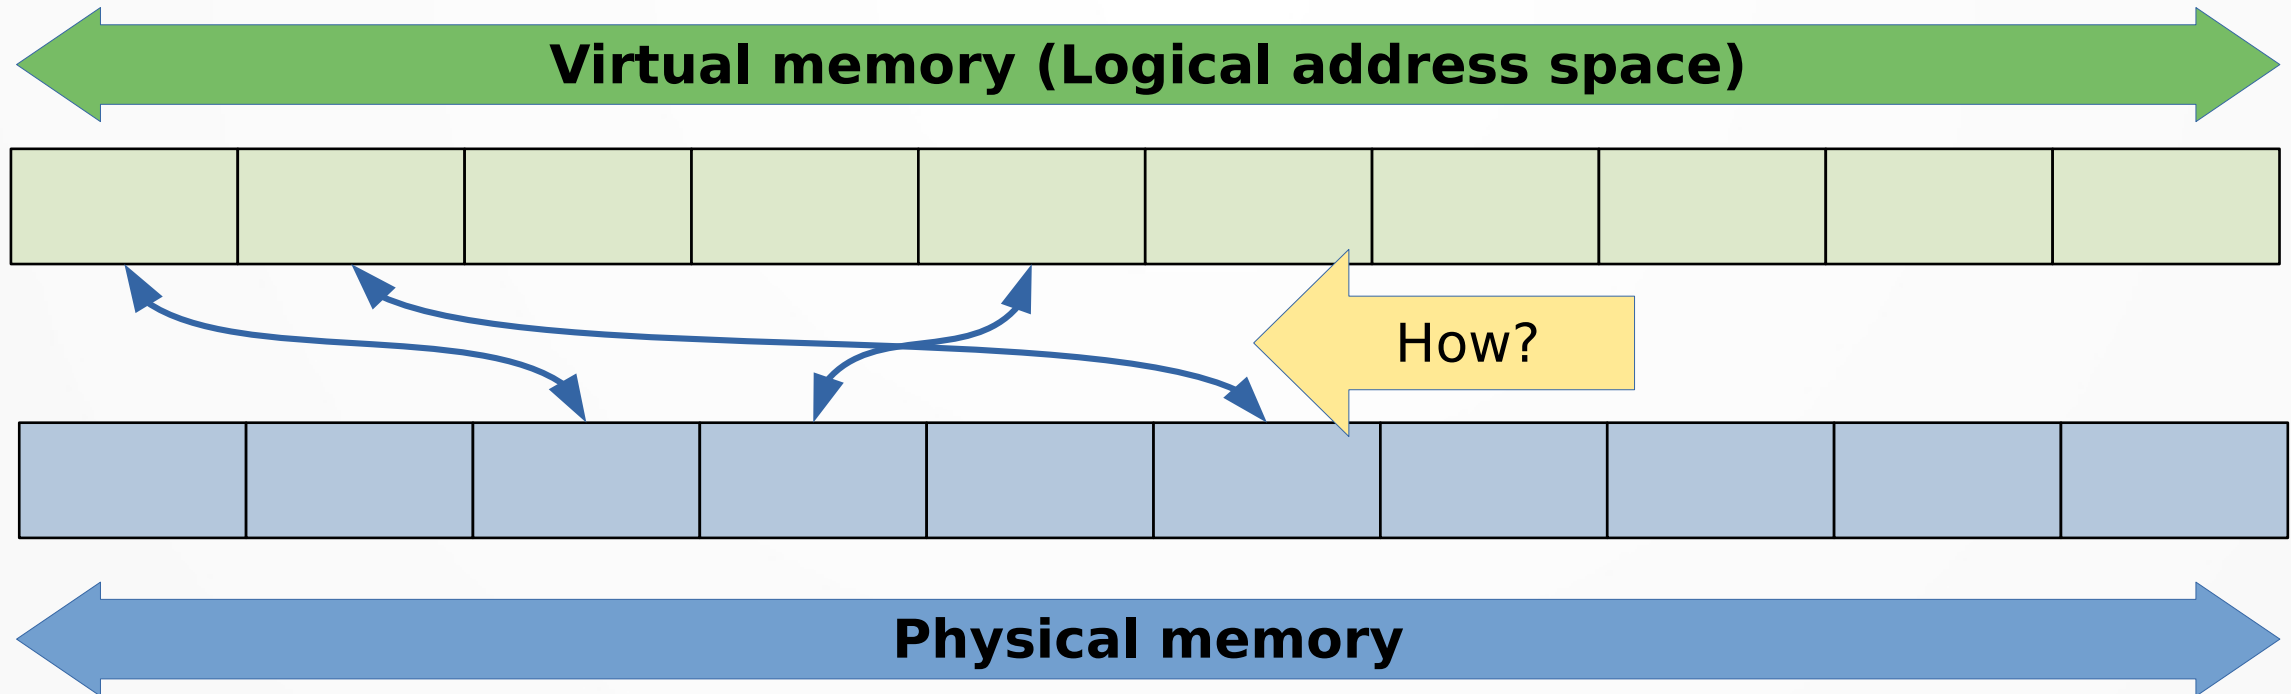
- Only pages that are used are mapped to physical pages
- Physical addresses are arbitrary and not visible to the program

# Virtual memory management



- Each process has its own virtual address space
- All virtual addresses are mapped to physical memory

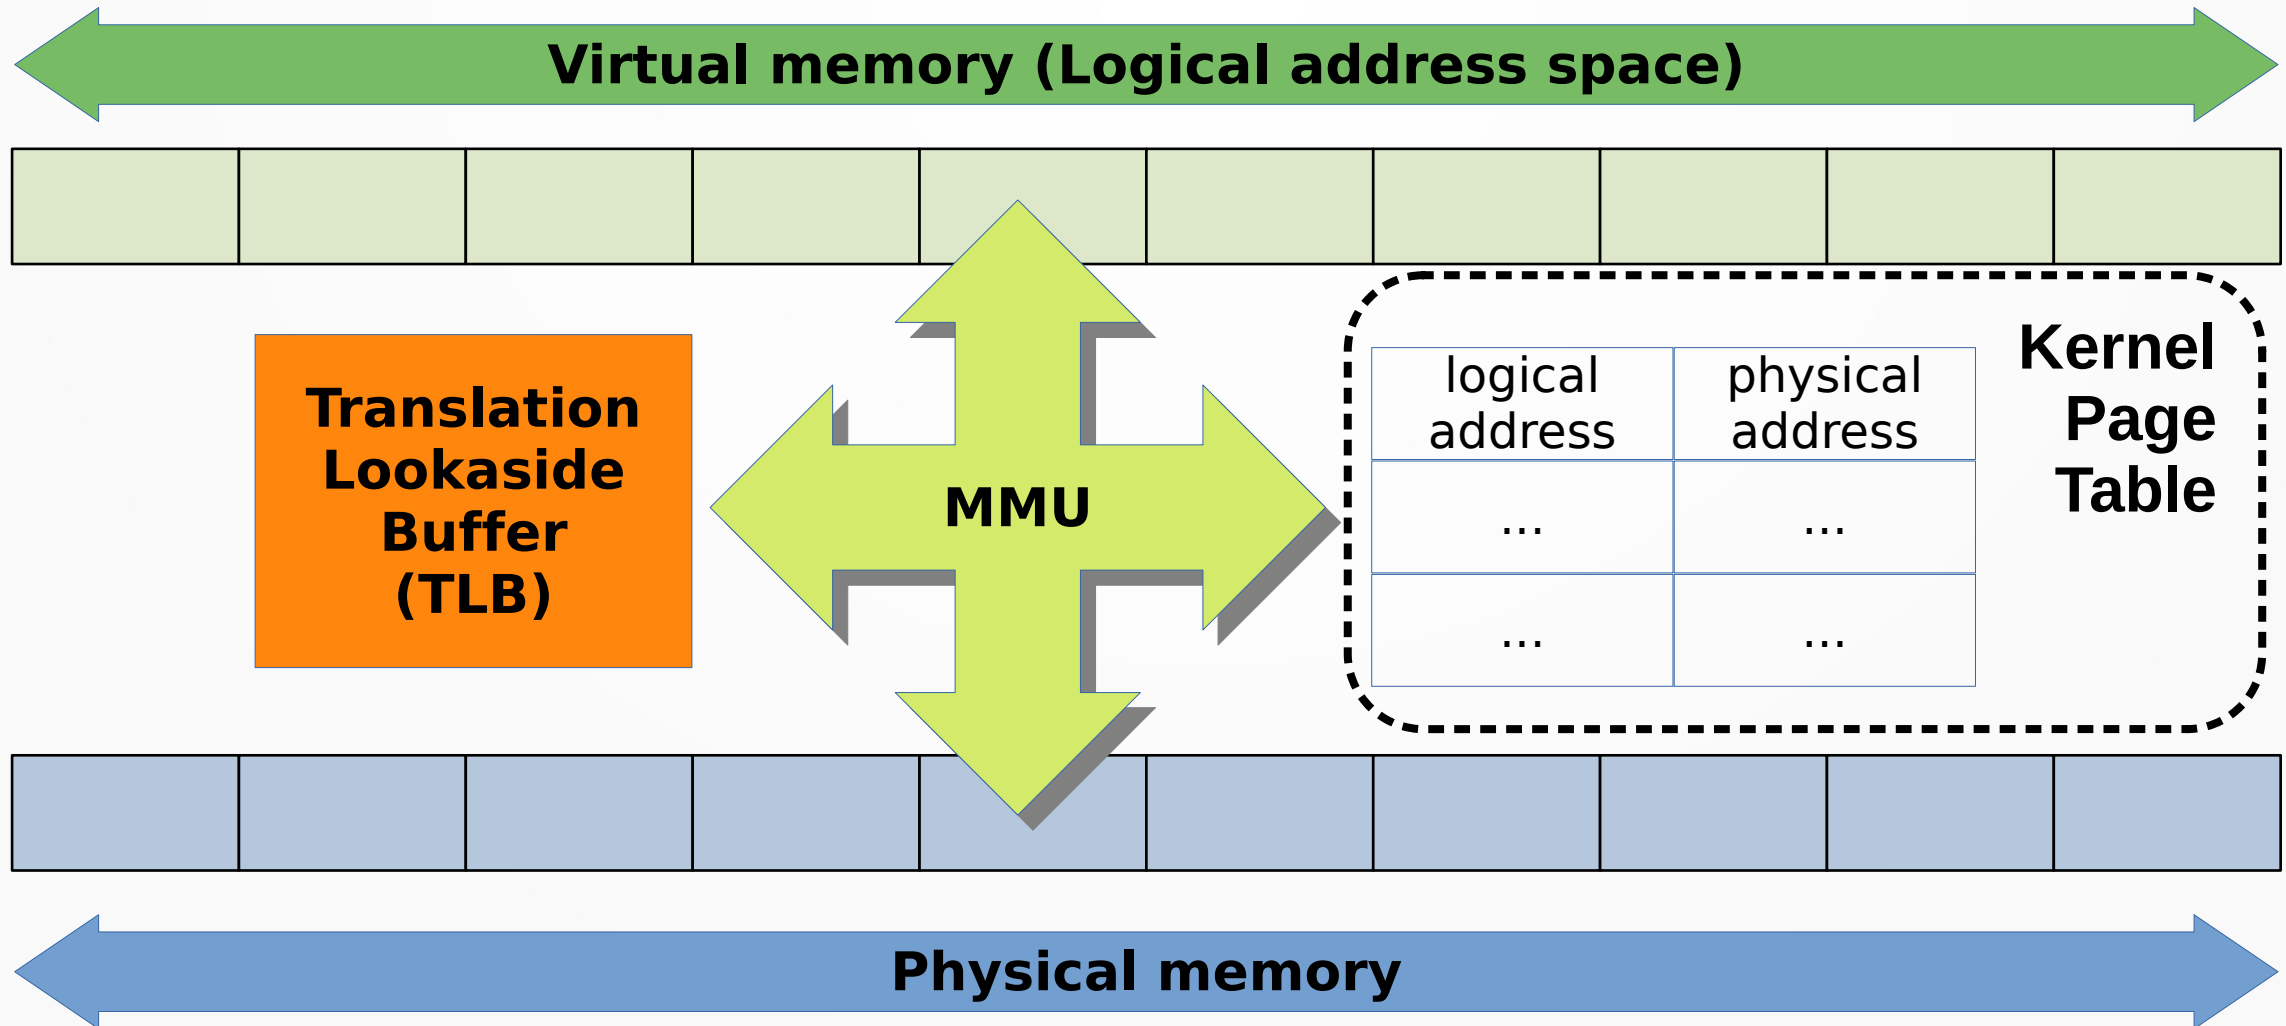
# Virtual memory management



- Mapping of logical to physical addresses is handled by the Memory Management Unit (MMU, part of the processor)
- MMU looks up the address map in the kernel page table



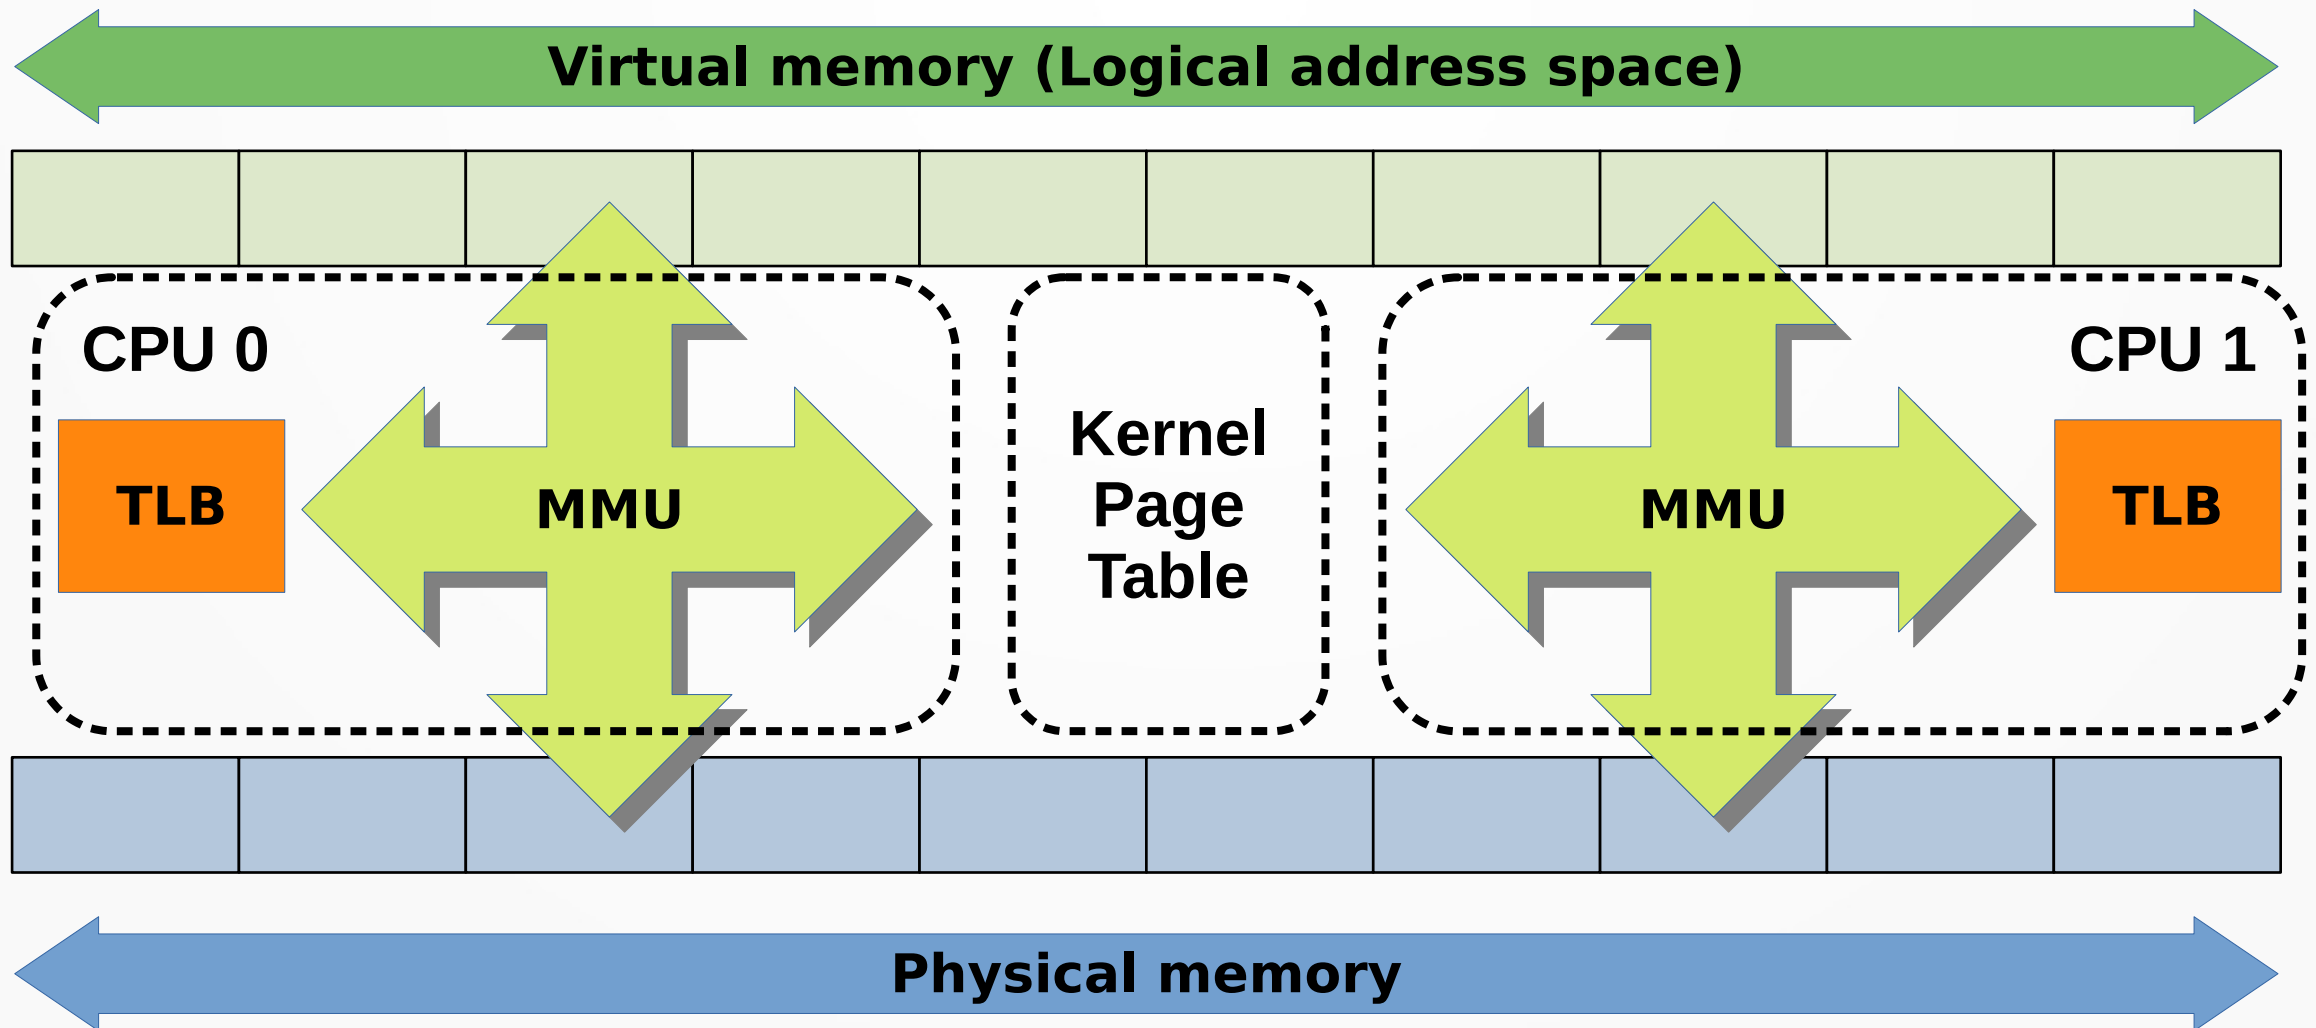
# Virtual memory management



# Virtual memory management

- The OS maintains the complete map of logical to physical addresses for all processes in the page table (per 4K page)
- The CPU has a hardware cache for recently used addresses
  - Translation lookaside buffer (TLB)
  - [Set-]Associative cache, similar to the regular CPU caches
    - Also known as Content-Addressable Memory
- If the address is found in the TLB (TLB hit), the MMU uses it
- If the address is not found in the TLB (TLB miss), the kernel has to walk the page table to find the map
- If the address is not found in the page table, page fault is triggered


# Virtual memory management in NUMA systems



# Virtual memory management in NUMA systems

- Logical address space is per process
  - Spans all NUMA nodes
- Physical address space is global
  - Each address belongs to a particular node but accessible by any CPU
- Logical to physical address map (page table) is global
- MMU and TLB are CPU circuits, generally one per node
  - Some architectures have several TLBs in one CPU – not important now
- TLB is similar to other caches: local cache of the global state
- Unlike CPU caches, there is no hardware TLB coherency
- TLB content can differ between CPUs

# Virtual memory management in NUMA systems



- If TLB content becomes outdated, the kernel must update it
  - There is no hardware assist for TLB updates

# Virtual memory management in NUMA systems

- TLB on different processors can get out of date
  - Unlike caches, there is no hardware TLB coherency
- The kernel updates the TLB using a “TLB shutdown”
  - *“Translation Lookaside Buffer Consistency: A Software Approach,” David L. Black, R. Rashid, D. Golub, C. R. Hill, R. Baron, ASPLOS 1989*
  - Update may affect just one address or many (possibly all)
  - TLB shutdown is an inter-processor interrupt
  - Kernel code runs on the CPU when a shutdown occurs
- TLB shutdown handling can be seen in the profiler (kernel calls) and is accounted as “system time”
- Why do shutdowns happen when that particular code is running?



# TLB shutdowns and their causes

- A TLB shutdown is processed by the kernel and interrupts the CPU
- TLB shutdown happens when a TLB is out of date, usually caused by
  - ~~Changes in memory mapping~~
  - ~~Memory access modes and protection~~
  - ~~Explicit changes to page table~~
  - NUMA migrations
- The profiler shows that the problem section spends its time in the kernel calls for TLB management and NUMA load balancing
- Migrating pages from one NUMA node to another changes the physical address and requires TLB updates

# Debugging NUMA and TLB shutdowns

- TLB shutdowns cannot be disabled but the profiler can show the time spent in the kernel processing them
  - Recent Intel CPUs have hardware counts for TLB flushes
- NUMA migration can be disabled
  - Bind the program to a single NUMA node
  - Turn off automatic NUMA balancing  
`echo 0 > /proc/sys/kernel/numa_balancing`
- Disabling NUMA migrations “solves” the performance slowdown
  - The overall program generally becomes slower
  - Auto balancing is a global state (affects everyone, needs root)

# Reducing the impact of TLB shootdowns

- TLB shootdowns cannot be disabled
- NUMA balancing shouldn't be disabled
- Why this particular code?
  - Happens in code that allocates and deallocates a lot of memory rapidly
    - The solution is to pre-allocate and reuse memory – we already do that
  - Happens in multi-threaded code that quickly walks over large memory
    - Changing this behavior may be possible but may incur costs elsewhere
- Can we make TLB shootdowns cheaper?
  - It's mainly the cost of the page table walk (one entry per 4kB)
  - Does it have to be 4kB?



# Page table optimization

- Page table problems are not new to NUMA
  - But can be worse in NUMA systems because of TLB inconsistency
- Often the solution is to enable “huge pages” – 2MB page
  - Must be enabled in the OS (usually on by default)  
`# cat /sys/kernel/mm/transparent_hugepage/enabled`  
`always [madvice] never`
  - `echo always > /sys/kernel/mm/transparent_hugepage/enabled`  
`echo madvice > /sys/kernel/mm/transparent_hugepage/enabled`
  - <https://www.kernel.org/doc/html/latest/admin-guide/mm/transhuge.html>
- “Madvice” means requested by the program
  - “Always” means the kernel can allocate huge pages as it sees fit
  - In practice, you cannot rely on that and must request huge pages

# Using huge pages in your program

- Converting a virtual address range to huge pages:  
`madvise(address, size, MADV_HUGEPAGE);`
  - address and size must be 4kB-aligned (size must be  $n \times 2\text{M}$  to be useful)
- The memory must come from `mmap()` private anonymous map
  - `mmap()` is used to reserve large chunks of virtual address space
    - No physical memory is reserved immediately
    - Page table may be optimized for large unmapped regions
    - Physical pages are reserved as page faults are processed
- Using huge pages reduces the number of page table/TLB entries
- Switching to huge pages resolved the observed slowdown



# Performance implications of NUMA

- What is NUMA?
- Accessing local memory is faster than accessing non-local memory
- Per-thread memory, data sharing, and cache locality are more important
- True data sharing can be significantly more expensive
- Simple benchmarks with different node binding often point to problems
- Hardware-assisted profiling is useful if you know what to look for
- Use good practices for memory access and node binding only if needed
- [Some] Concurrent data structures need to be redesigned for NUMA
- For I/O-bound programs, use the right node for the data
- There's going to be something you don't know about...



# Nearly Unfathomable Morass of Arcana (NUMA)

Questions?

Possibly answers too...